

# CSDS503 / COMP552 – Advanced Machine Learning

---

Faizad Ullah

# Nearest Neighbors Methods

---

# The K Nearest Neighbors Algorithm


- Every Machine Learning Algorithm comes up with certain assumptions.
- “A data point is known by the company it keeps” (Aesop – the data scientist)
- “A data point is known by the most common company it keeps” (Aesop – the intelligent data scientist)
- A vector is known by the company it keeps

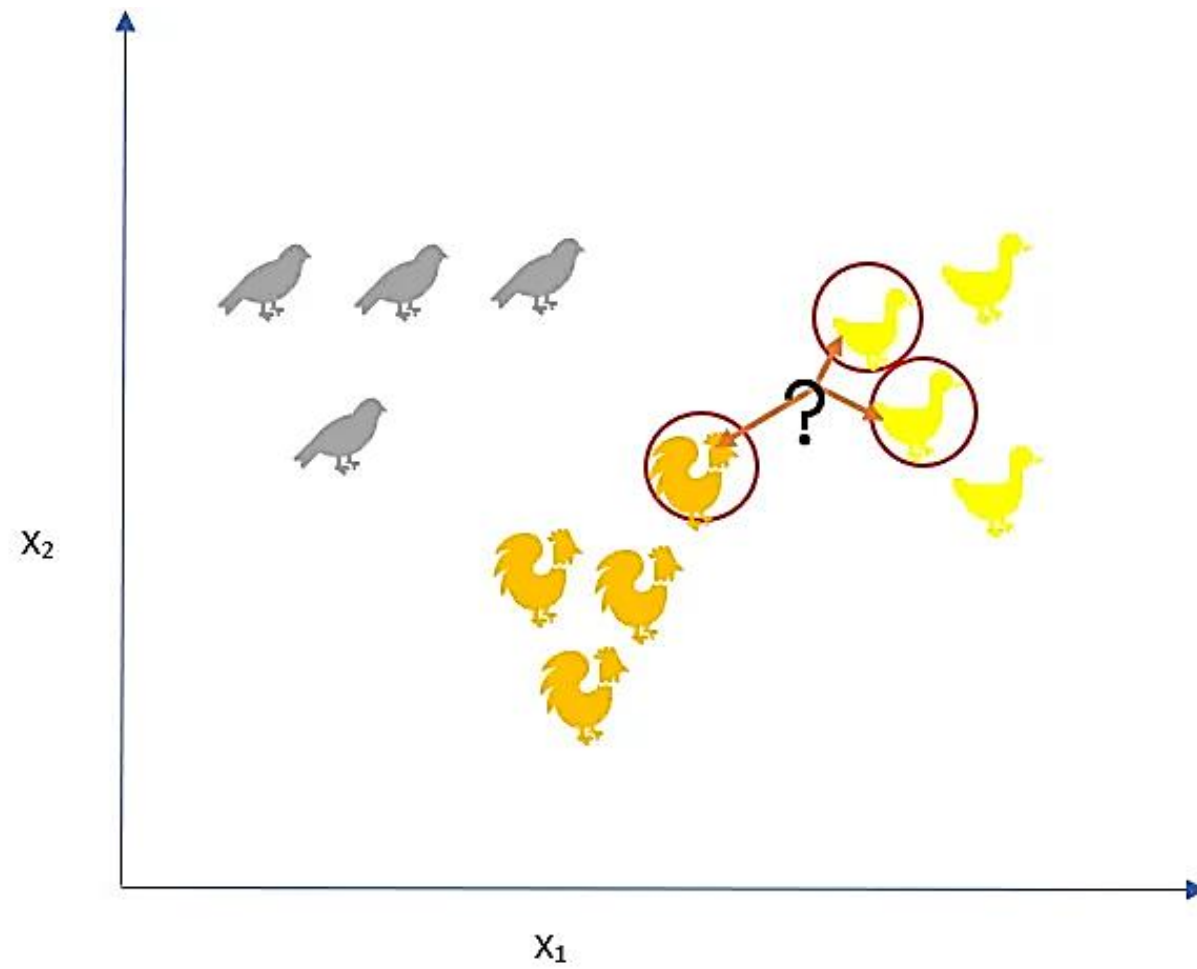
# The K Nearest Neighbors Algorithm

- KNN relies on the concepts of proximity and similarity.
- KNN is a supervised learning algorithm capable of performing both classification and regression tasks.

# The K Nearest Neighbors Algorithm

- **Basic idea:** Similar **Inputs** have similar **outputs**
- **Classification rule:** For a **test input  $x$** , assign the **most common label** amongst its  **$k$  most similar** (nearest) training inputs.





# Formal Definition

- Assuming  $x$  to be our test point, let's denote the set of the  $k$  nearest neighbors of  $x$  as  $S_x$
- Formally,  $S_x$  is defined as:

$$\begin{aligned} S_x &\subseteq D \text{ s.t. } |S_x| = k \\ &\text{and} \\ \forall (x', y') \in D \setminus S_x, \\ \text{dist}(x, x') &\geq \max_{(x'', y'') \in S_x} \text{dist}(x, x'') \end{aligned}$$

Every point that is in  $D$  but not in  $S_x$  is at least as far away from  $x$  as the furthest point in  $S_x$

# Formal Definition



- We define the classifier  $h()$  as a function returning the most common label in  $S_x$ :

$$h(x) = \text{mode}(\{y'' : (x'', y'') \in S_x\})$$

- Where  $\text{mode}(\cdot)$  means to select the label of the highest occurrence.
- So, what do we do if there is a draw?
- Keep  $k$  odd (for binary classification) or return the result of  $(k-1)$ NN with a smaller  $k$ .



# KNN Decision Boundary



# Properties of KNN – Non-parametric

The KNN Algorithm is a supervised, nonparametric algorithm.

It does not make any assumptions about the underlying distribution nor tries to estimate it.

# Properties of KNN – Non-parametric

**Parametric models** summarize data with a fixed set of parameters (independent of the number of training examples).

No matter how much data you throw at a parametric model, it won't change its mind about how many parameters it needs.

# Properties of KNN – Non-parametric

**Non-parametric** models make no assumptions about the probability distribution or number of parameters when modeling the data.

**Non-parametric** methods are good when you have a lot of data and no prior knowledge, and when you don't want to worry too much about choosing just the right features.

# Properties of KNN – Non-parametric

**Non-parametric** does not mean that they have **no parameters!**

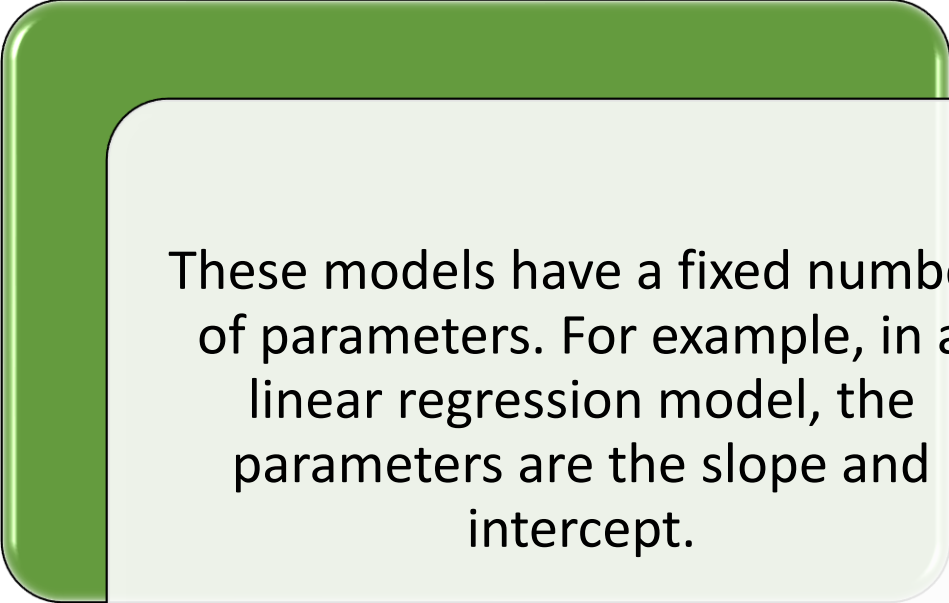
On the contrary, **non-parametric** models (can) become more and more complex with an increasing amount of data.

# Properties of KNN – Non-parametric

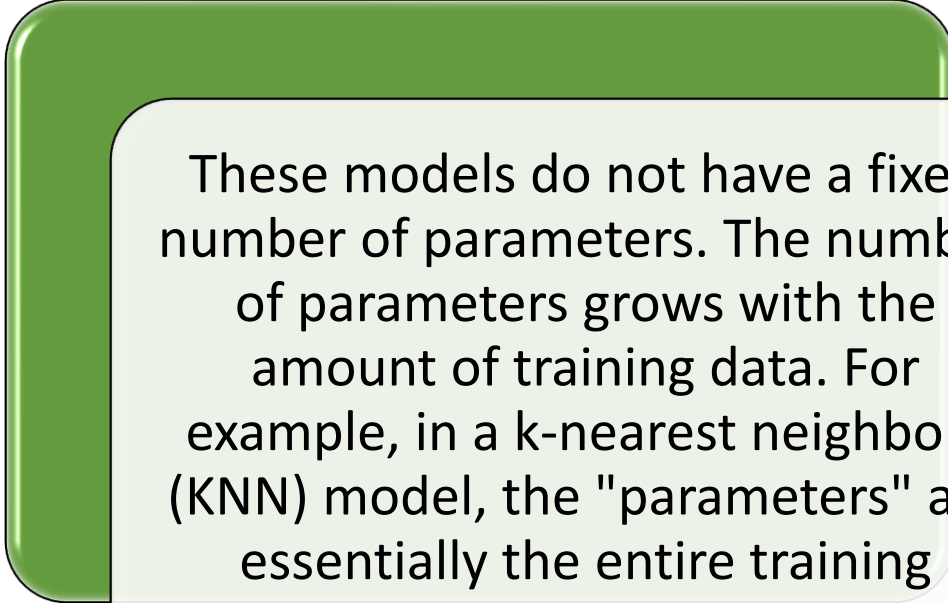
In a **parametric model**, we have a **finite** number of **parameters**, and in **non-parametric** models, the number of parameters is (potentially) **infinite**.

Still, the distinction is a bit ambiguous at best.

# Parametric vs. Non-Parametric Models

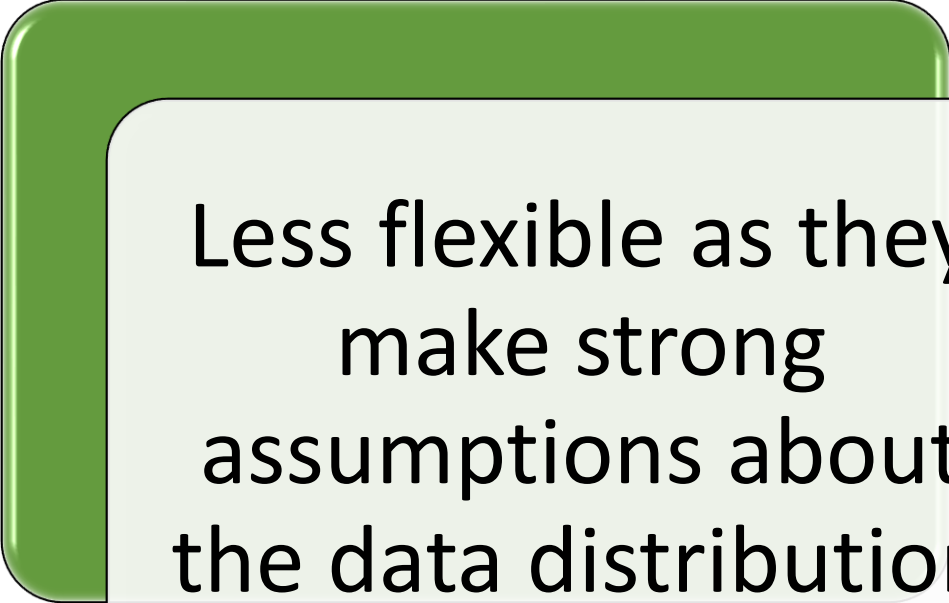


These models have a fixed number of parameters. For example, in a linear regression model, the parameters are the slope and intercept.

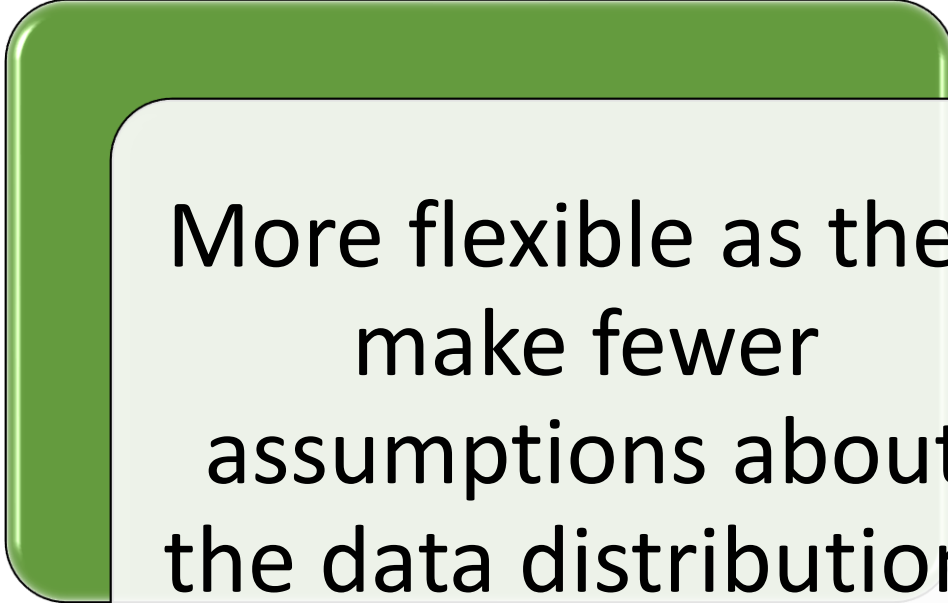


These models do not have a fixed number of parameters. The number of parameters grows with the amount of training data. For example, in a k-nearest neighbors (KNN) model, the "parameters" are essentially the entire training dataset.

# Parametric vs. Non-Parametric Models



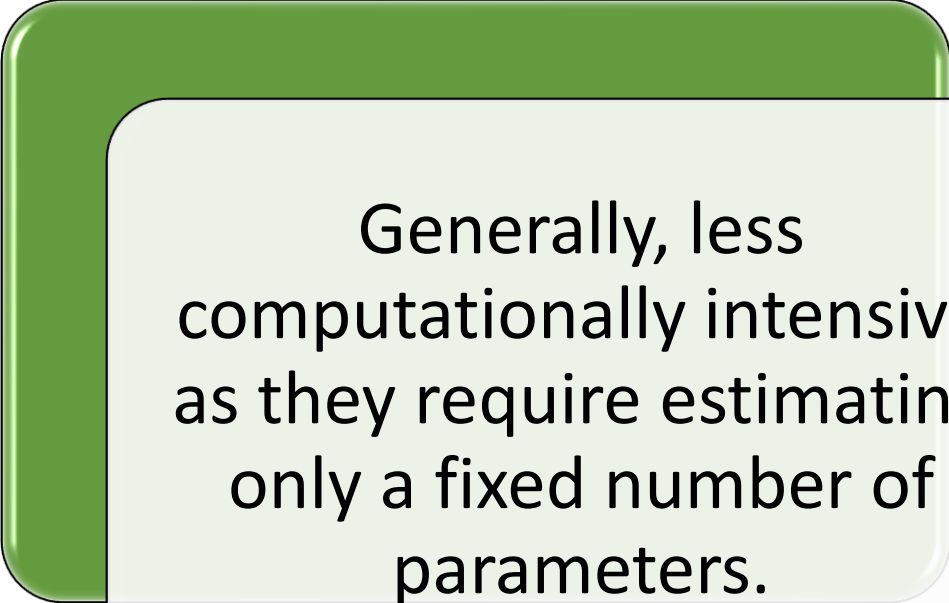
Less flexible as they  
make strong  
assumptions about  
the data distribution



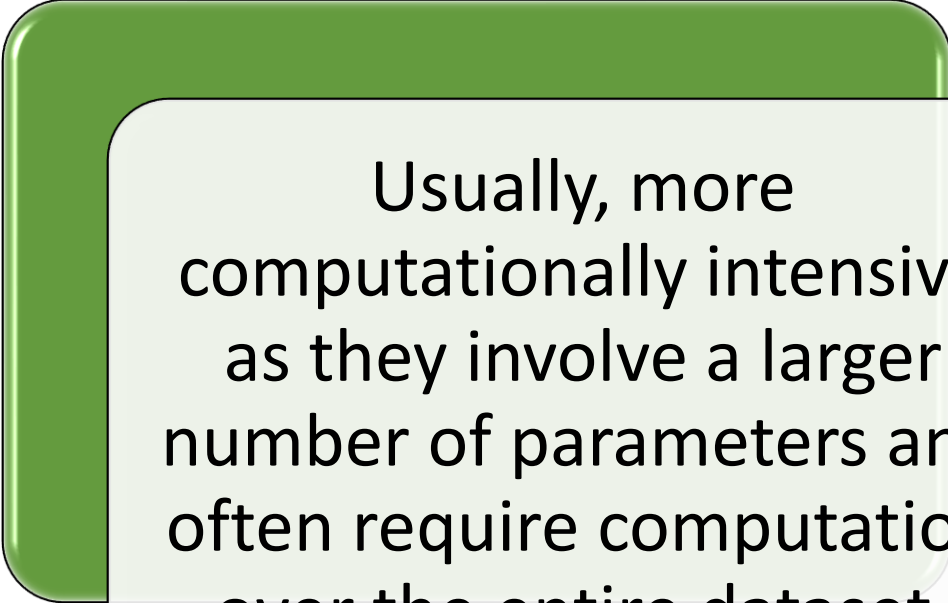
More flexible as they  
make fewer  
assumptions about  
the data distribution.



# Parametric vs. Non-Parametric Models

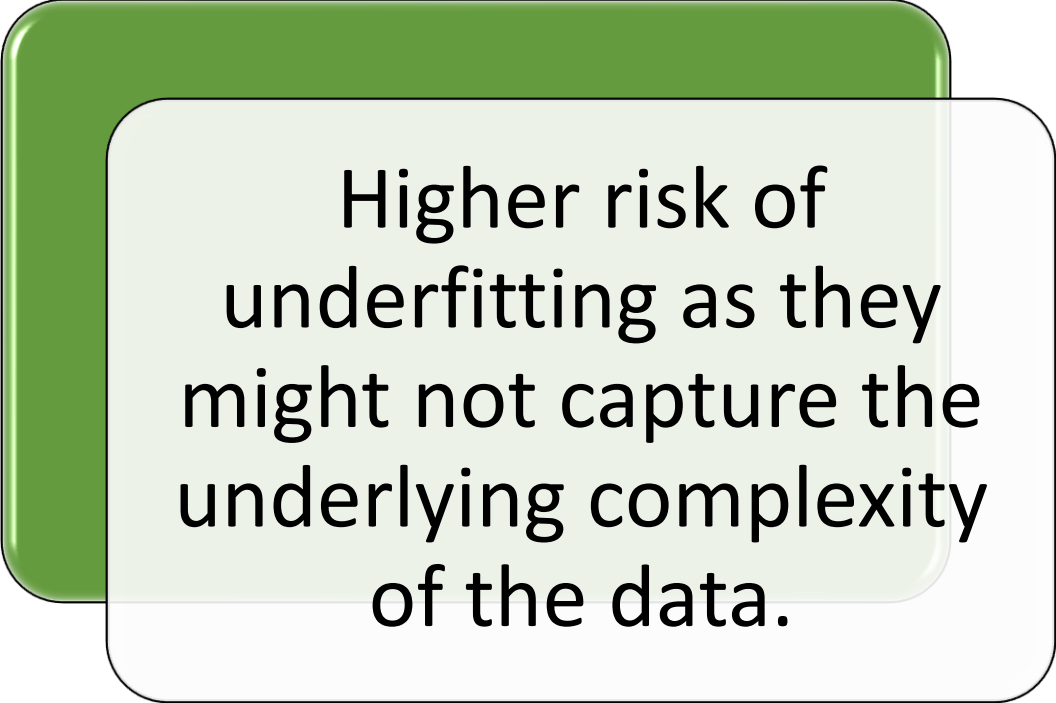


Generally, less computationally intensive as they require estimating only a fixed number of parameters.



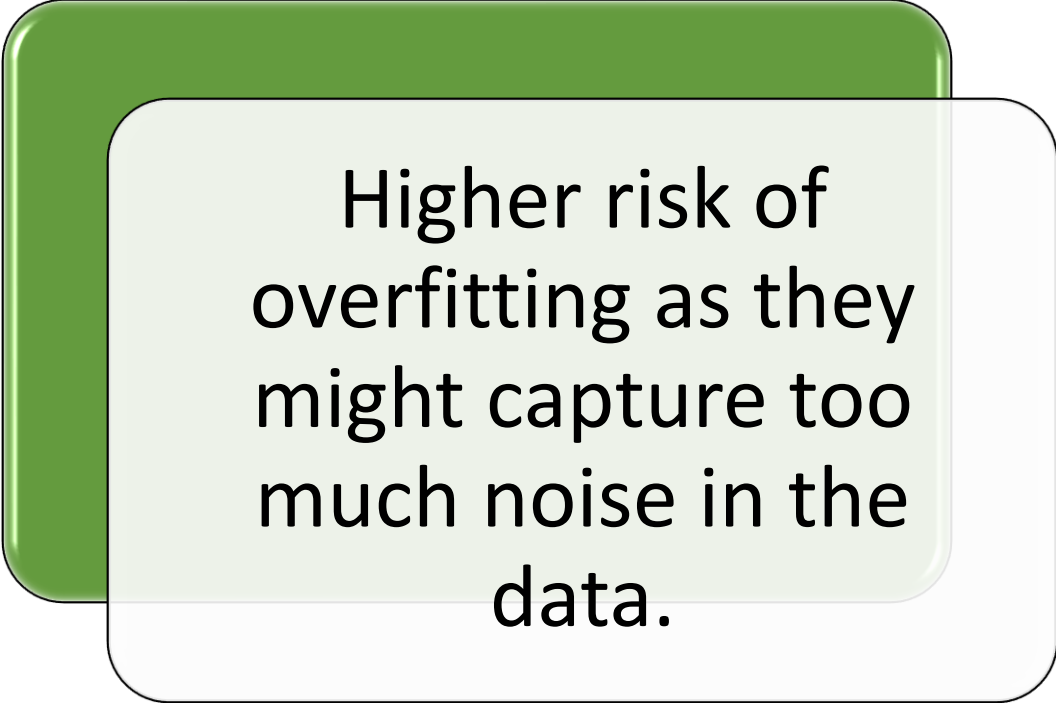
Usually, more computationally intensive as they involve a larger number of parameters and often require computation over the entire dataset.

# Parametric vs. Non-Parametric Models



Higher risk of underfitting as they might not capture the underlying complexity of the data.

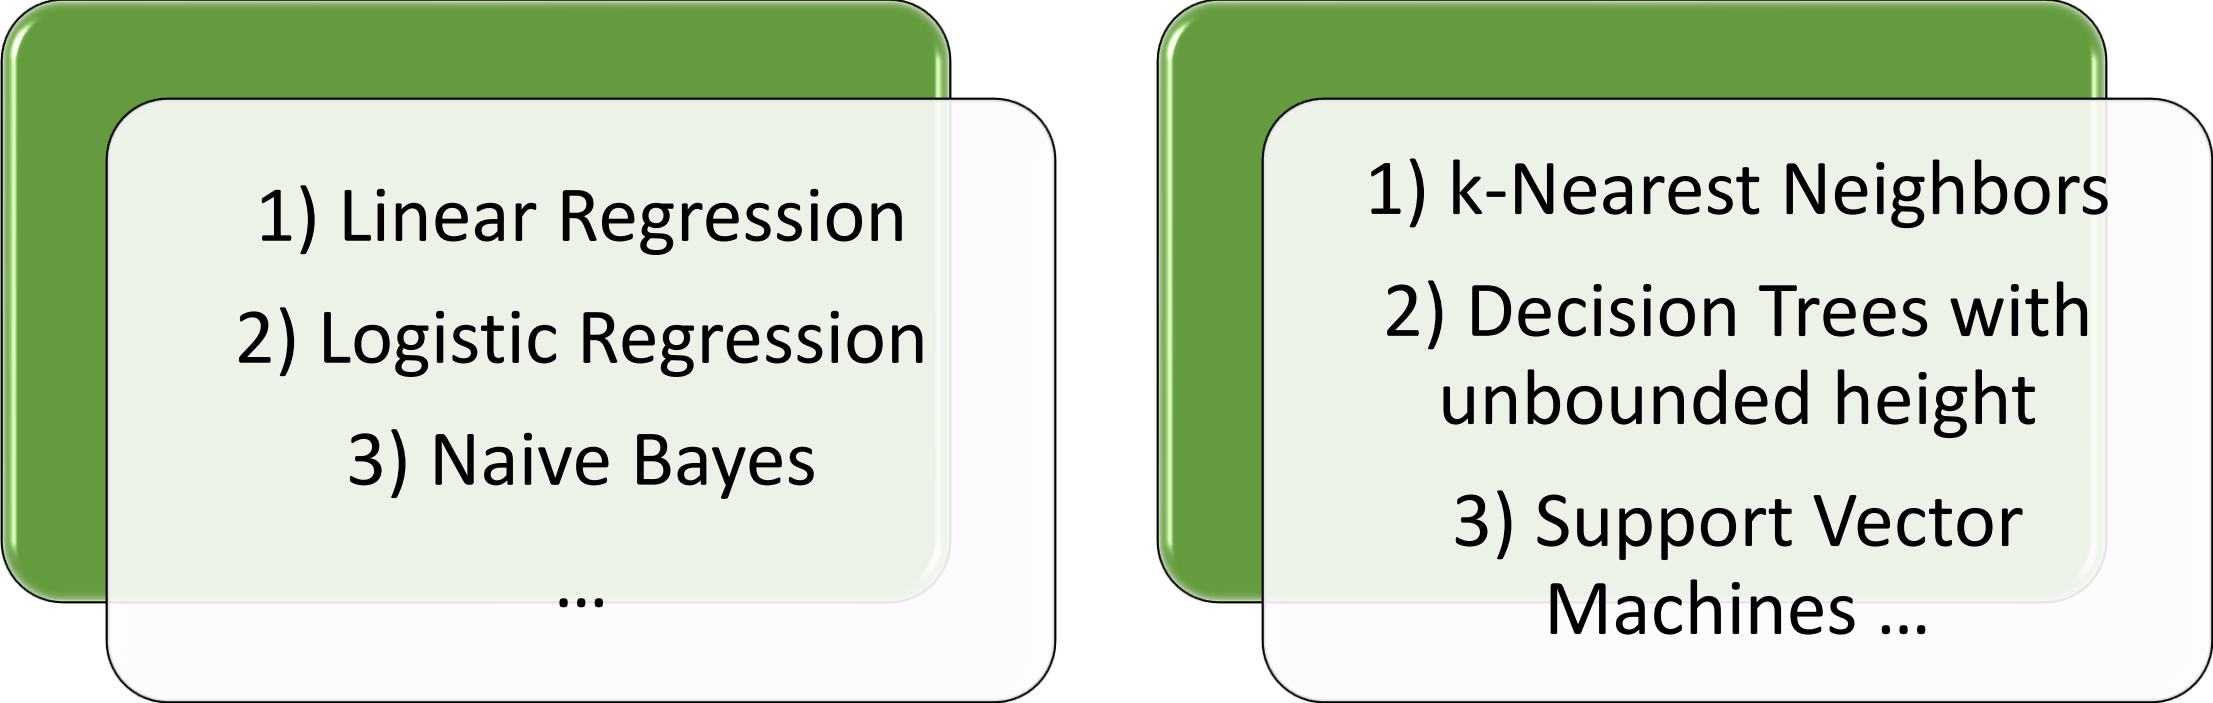
The diagram consists of a light green rounded rectangle with a dark green header. Inside, a white rounded rectangle contains the text. The white rectangle is slightly offset to the right and bottom, creating a layered effect.



Higher risk of overfitting as they might capture too much noise in the data.

The diagram consists of a light green rounded rectangle with a dark green header. Inside, a white rounded rectangle contains the text. The white rectangle is slightly offset to the right and bottom, creating a layered effect.

# Parametric vs. Non-Parametric Models

- 
- 1) Linear Regression
  - 2) Logistic Regression
  - 3) Naive Bayes

...

- 1) k-Nearest Neighbors
- 2) Decision Trees with unbounded height
- 3) Support Vector Machines ...

# Properties of KNN – Non-parametric

- **Parameters** and **hyperparameters** are two types of values that a model uses to make predictions, but they serve different purposes and are learned in different ways:
- **Parameters:**
  - These are the parts of the model that are learned from the training data during the training process.
  - For example, the weights in a linear regression model are parameters.
  - The model uses the training data to adjust the parameters to minimize the prediction error.

# Properties of KNN – Non-parametric

- **Hyperparameters:**

- These are the settings or configurations that need to be specified prior to training the model.
- They are not learned from the data but are essential for the learning process.
- For example, the learning rate in gradient descent, the depth of a decision tree, or the number of clusters in k-means clustering are all hyperparameters.
- The values of hyperparameters are usually set before training the model and remain constant during the training process. They may be adjusted between runs of training to optimize model performance.

# Properties of KNN – Non-parametric

- **Classification:** Choose the most frequent class label amongst k-nearest neighbors
- **Regression:** Take an average over the output values of the k-nearest neighbors and assign to the test point.

# Properties of KNN – Non-parametric

- An Instance-based learning algorithm

- Instead of performing explicit generalization, form hypotheses by comparing new problem instances with training instances
- (+) Can easily adapt to unseen data
- (-) Complexity of prediction is a function of  $n$  (size of training data)

- A lazy learning algorithm

- Delay computations on training data until a query is made, as opposed to eager learning
- (+) Good for continuously updated training data like recommender systems
- (-) Slower to evaluate and need to store the whole training data

# Distance/Similarity

---



# Similarity/Distance Measures

- If scaled between 0 and 1, then  $sim = 1 - distance$
- The Minkowski distance is a generalized metric form of Euclidean, Manhattan and Chebyshev distances

# Minkowski Distance

- The Minkowski distance between two n-dimensional
- vectors  $P = \langle p_1, p_2, \dots, p_n \rangle$  and  $Q = \langle q_1, q_2, \dots, q_n \rangle$ , it is
- defined as:

$$d_{minkowski}(p, q) = \left( \sum_{i=1}^n |p_i - q_i|^a \right)^{1/a}, a \geq 1$$

- $a = 1$ , is the Manhattan distance
- $a = 2$ , is the Euclidean distance
- $a \rightarrow \infty$ , is the Chebyshev distance

# Constraints on Distance Metrics

- The distance function between vectors  $p$  and  $q$  is a function  $d(p, q)$  that defines the distance between both vectors is considered as a metric if it satisfy a certain number of properties:

1. **Non-negativity:** The distance between  $p$  and  $q$  is always a value greater than or equal to zero

- $d(p, q) \geq 0$

2. **Identity of indiscernible vectors:** The distance between  $p$  and  $q$  is equal to zero if and only if  $p$  is equal to  $q$

- $d(p, q) = 0 \text{ iff } p = q$

# Constraints on Distance Metrics



3. **Symmetry**: The distance between  $p$  and  $q$  is equal to the distance between  $q$  and  $p$ .

$$d(p, q) = d(q, p)$$

4. **Triangle inequality**: Given a third point  $r$ , the distance between  $p$  and  $q$  is always less than or equal to the sum of the distance between  $p$  and  $r$  and the distance between  $r$  and  $q$

$$d(p, q) \leq d(p, r) + d(r, q)$$


# Euclidean vs Manhattan vs Hamming Distance



Data A: 1 1 1 0 0 1 0 1 0 1 0 1 1 0 0 0 1 0 0 1 0 0 1 1 1 0 0

Data B: 1 1 1 0 0 1 0 1 0 1 0 1 1 0 1 0 0 1 0 0 1 0 0 1 1 1 0 0

Hamming Distance between Data A and  
Data B = 1



Taxicab geometry versus  
Euclidean distance: In taxicab  
geometry, the red, yellow, and blue  
paths all have the same shortest  
path length of 12. In Euclidean  
geometry, the green line has length  
 $6\sqrt{2} \approx 8.49$ , and is the unique  
shortest path.

# Manhattan Distance

- Also known as Manhattan length, rectilinear distance, L1 distance or L1 norm, city block distance, snake distance, taxi-cab metric, or city block distance

$$d_{Man}(p, q) = d(q, p) = |p_1 - q_1| + |p_2 - q_2| + \dots + |p_n - q_n|$$
$$d(p, q) = d(q, p) = \sum_{i=1}^n |p_i - q_i|$$

# Euclidean Distance

$$d(p, q) = d(q, p) = \sqrt{(p_1 - q_1)^2 + (p_2 - q_2)^2, \dots, (p_n - q_n)^2}$$

$$d(p, q) = d(q, p) = \sqrt{\sum_{i=1}^n (p_i - q_i)^2}$$

# Euclidean Distance

- Good choice for numeric attributes
- When data is dense or continuous, this is a good proximity measure
- Downside: Sensitive to extreme deviations in attributes (as it squares differences)
- The variables which have the largest value greatly influence the result
- Does not work well for situations where features on different scales are mixed (e.g., #bedrooms (1-5) and area (200 – 5,000 sq feet) of a house)
- Solution: feature normalization (min-max scaling)



# Chebyshev Distance

- For Chebyshev distance, the distance between two vectors is the greatest of their differences along any coordinate dimension
- When two objects are to be defined as “different”, if they are different in any one dimension
- Also called chessboard distance, maximum metric, or  $L^\infty$  metric

$$d_{cheb}(p, q) = \max_i |p_i - q_i|$$

# Chebyshev Distance

$$d_{Cheb}(p, q) = \lim_{a \rightarrow \infty} \left( \sum_{i=1}^n |p_i - q_i|^a \right)^{1/a} = \max_i |p_i - q_i|$$

How?

Assume,  $p = \langle 2, 3, \dots, 9 \rangle$ ,  $q = \langle 4, 6, \dots, 10 \rangle$

$$d_{Cheb}(p, q) = \lim_{a \rightarrow \infty} (|2 - 4|^a + |3 - 6|^a + \dots + |9 - 10|^a)^{\frac{1}{a}}$$

$$d_{Cheb}(p, q) = \lim_{a \rightarrow \infty} (2^a + 3^a + \dots + 1^a)^{\frac{1}{a}}$$

Suppose,  $a = 2$

$$d(p, q) = (4 + 9 + \dots + 1)^{\frac{1}{2}}$$

Suppose,  $a = 3$

$$d(p, q) = (8 + 27 + \dots + 1)^{\frac{1}{3}}$$

Suppose,  $a = 10$

$$d(p, q) = (1,024 + 59,049 + \dots + 1)^{\frac{1}{10}}$$

Now,  $a \rightarrow \infty$

$$d_{Cheb}(p, q) = \lim_{a \rightarrow \infty} \left( \sum_{i=1}^n |p_i - q_i|^a \right)^{\frac{1}{a}} \rightarrow \max_i |p_i - q_i|^a)^{\frac{1}{a}}$$

$$d_{Cheb}(p, q) = \lim_{a \rightarrow \infty} (\max_i |p_i - q_i|^a)^{\frac{1}{a}}$$

# The KNN Algorithm

---

# The KNN Algorithm

- **Input:**

- Training samples  $D = \{ (x_1, y_1), (x_2, y_2), (x_3, y_3), \dots, (x_n, y_n) \}$
- Test sample  $d = (x, y)$
- Assume  $x$  to be an  $m$ -dimensional vector

- **Output:** Class label of test sample  $d$

1. Compute the distance between  $d$  and every sample in  $D$
2. Choose the  $K$  samples in  $D$  that are nearest to  $d$ ; denote the set by  $S_d \in D$
3. Assign  $d$  the label  $y_i$  of the majority class in  $S_d$

- **Note:**


- All action takes place in the test phase; the training phase is essentially to clean, normalize, and store the data

# The KNN Algorithm


#	Height (inches)	Weight (kgs)	B.P. Sys	B.P. Dia	Heart disease	Cholesterol Level	Euclidean Distance
1	62	70	120	80	No	150	52.59
2	72	90	110	70	No	160	47.81
3	74	80	130	70	No	130	43.75
4	65	120	150	90	Yes	200	7.14
5	67	100	140	85	Yes	190	16.61
6	64	110	130	90	No	130	15.94
7	69	150	170	100	Yes	250	44.26
8	66	115	145	90			

# The KNN Algorithm

- Handwritten digit recognition
  - 16x16 bitmaps
  - 8-bit grayscale
  - Euclidean distances over raw pixels



# The KNN Algorithm



- **Accuracy:**
  - 7-NN ~ 95.2%
  - SVM ~ 95.8%
  - Humans ~ 97.5%


$$D(x, y) = \sqrt{\sum_{i=0}^{255} (x_i - y_i)^2}$$

# The KNN Algorithm

## 1. Text Classification using KNN:

1. Machine learning is fascinating.
2. NLP is essential.
3. Deep learning is interesting.
4. Text classification is essential.
5. Data science is interesting.

## 2. Image Classification using KNN:





# The KNN Algorithm

Document	machine	learning	is	fascinating	NLP	essential	deep	interesting	text	classification	data	science	Label
1	1	1	1	1	0	0	0	0	0	0	0	0	ML
2	0	0	1	0	1	1	0	0	0	0	0	0	TEXT
3	0	1	1	0	0	0	1	1	0	0	0	0	ML
4	0	0	1	0	0	1	0	0	1	1	0	0	TEXT
5	0	0	1	0	0	0	0	1	0	0	1	1	?

# Time Complexity of KNN

---

# Complexity of KNN

- Input:
  - Training samples  $D = \{ (x_1, y_1), (x_2, y_2), \dots, (x_n, y_n) \}$
  - Test sample  $d = (x_1, y_1)$ ,  $k$ . Assume  $x$  to be an  $m$ -dimensional vector.
- Output: Class label of test sample  $d$ 
  1. Compute the distance between  $d$  and every sample in  $D$ ;  $n$  samples, each is  $m$ -dimensional  $\Rightarrow O(mn)$
  2. Choose the  $K$  samples in  $D$  that are nearest to  $d$ ; denote the set by  $S_d \in D$ 
    - Either naively do  $K$  passes of all samples costing  $O(n)$  each time for  $O(nk)$
    - Or use the quickselect algorithm to find the  $k$ th smallest distance in  $O(n)$  and then return all distances no larger than the  $k$ th smallest distance. This will accumulate to  $O(n)$
  3. Assign  $d$  the label  $y_i$  of the majority class in  $S_d$ , this is  $O(k)$ .

# Complexity of KNN

- Time complexity:

$$O(mn + n + k) = O(mn)$$

- Space complexity:

$$O(mn)$$

# Tuning the Hyperparameter K

- Divide your training data into training and validation sets.
- Do multiple iterations of **m-fold cross-validation**, each time with a different value of k, starting from k=1
- Keep iterating until the k with the best classification accuracy (minimal loss) is found

# Tuning the Hyperparameter K

- What happens if we use the training set itself as the test dataset instead of a validation set? Which k wins?
- $K=1$ :
  - As there is always the nearest instance with the correct label: The instance itself!
- $K=n$ :
  - KNN will always return the majority class in the dataset

# Tuning the Hyperparameter K

- KNN is a simple algorithm but is highly effective for solving various real-life classification problems.
- Especially when the datasets are large and continuously growing.
- **Challenges:**
  - How to find the optimum value of K?
  - How to find the right distance function?
- **Problems:**
  1. High computational time cost for each prediction.
  2. High memory requirement as we need to keep all training samples.
  3. The curse of dimensionality.