

# COMP 554 / CSDS 553 Advanced NLP

---

Faizad Ullah

# Basics of Neural Networks

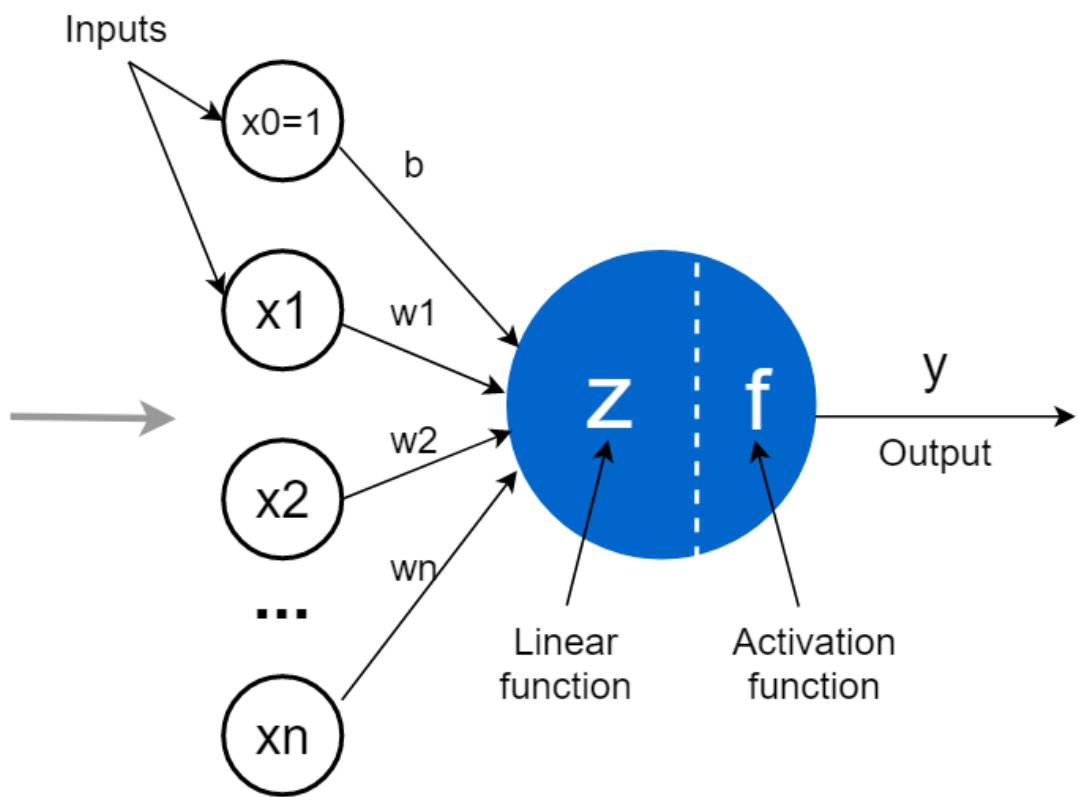
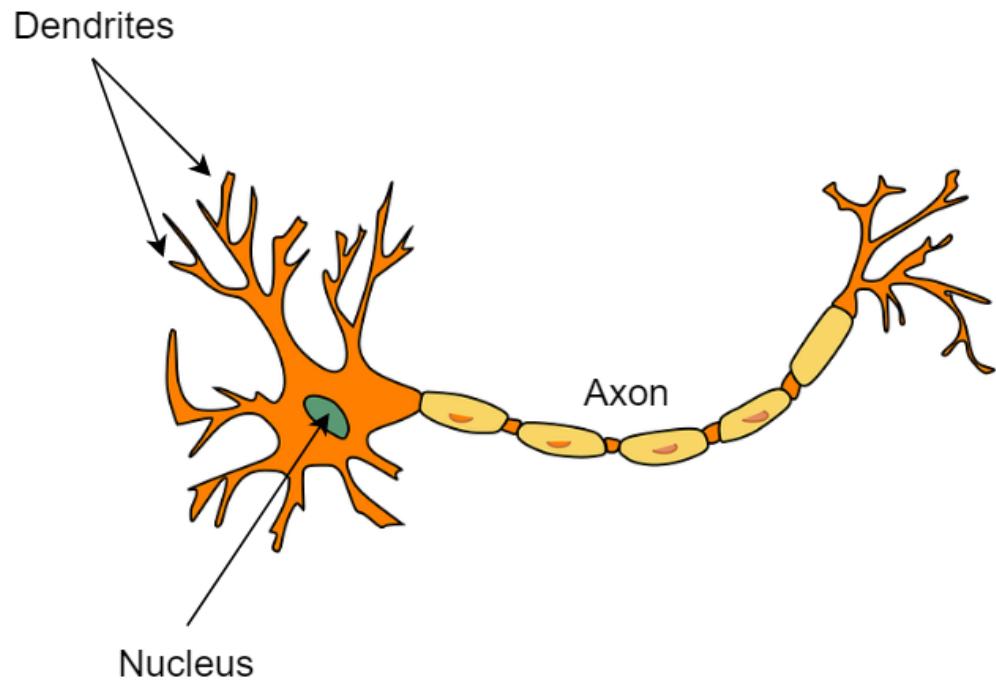
---

# Four Components for ML Systems

- A feature representation of the input, i.e.,  $[x_1, x_2, \dots, x_n]$
- A classification function that computes  $\hat{y}$ , the estimated class, via  $p(y|x)$ . Sigmoid, softmax, etc.
- An objective function that we want to optimize for learning, usually involving minimizing a loss function corresponding to error on training examples. Cross-entropy loss function, MAE, MSE, etc.
- An algorithm for optimizing the objective function. The gradient descent

# Four Components for ML Systems

- **Training:** We train the system (specifically the weights  $w$  and  $b$ ) using stochastic gradient descent and the cross-entropy loss.
- **Testing:** Given a test example  $x$  we compute  $p(y|x)$  and return the higher probability label  $y = 1$  or  $y = 0$ .

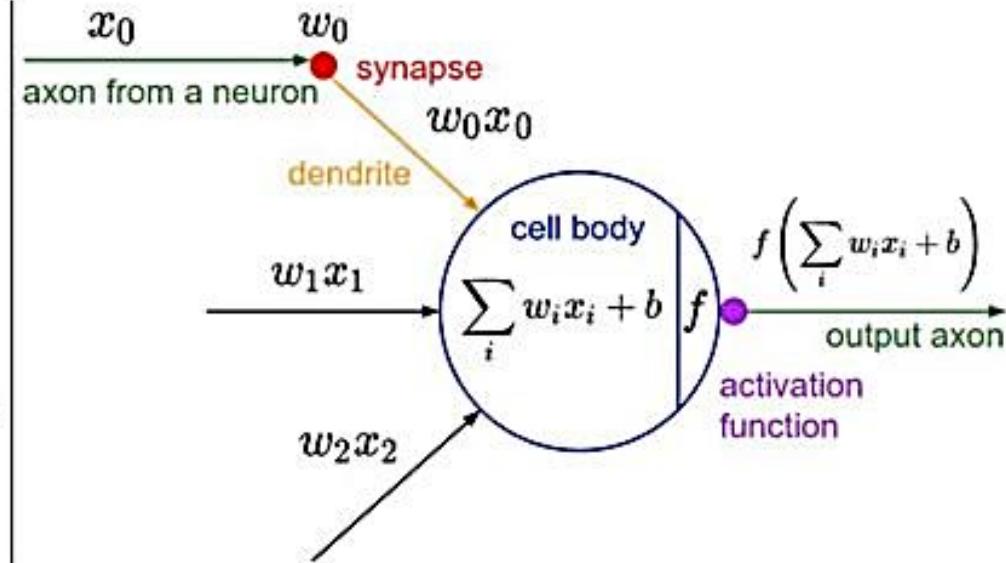
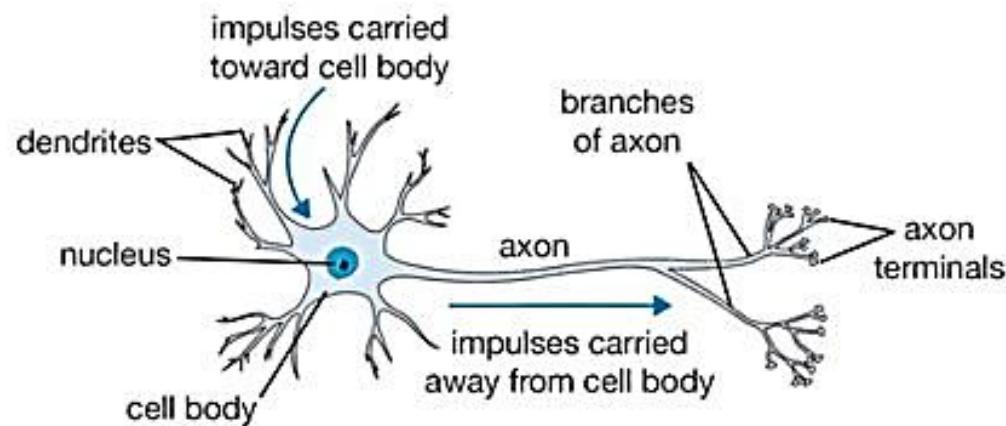


**Dendrites** receive signals from other neurons

**Soma** processes the information

**Axons** transmit the output

**Synapses** are the connections to other neurons



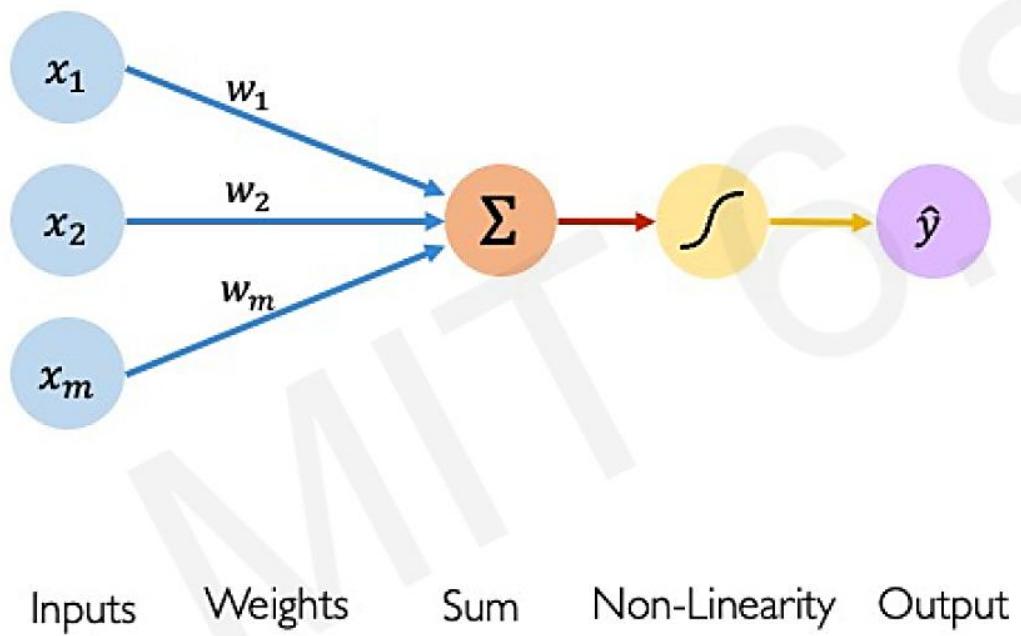
**Dendrites** receive signals from other neurons

**Soma** processes the information

**Axons** transmit the output

**Synapses** are the connections to other neurons

# Perceptron



Output

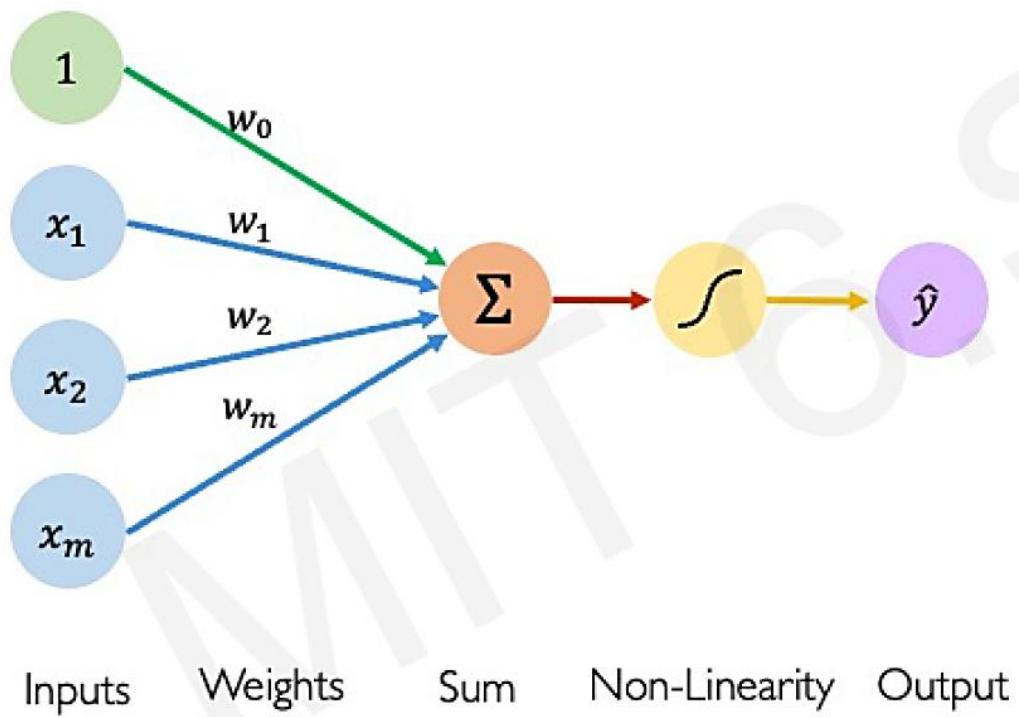
Linear combination of inputs

$\hat{y} = g \left( \sum_{i=1}^m x_i w_i \right)$

Non-linear activation function

The equation  $\hat{y} = g \left( \sum_{i=1}^m x_i w_i \right)$  is shown with annotations explaining its components. A red arrow points to the summation term  $\sum_{i=1}^m x_i w_i$  and is labeled "Linear combination of inputs". A yellow arrow points to the activation function  $g$  and is labeled "Non-linear activation function". A purple arrow points to the final output  $\hat{y}$ .

# Perceptron



Linear combination of inputs  $\downarrow$

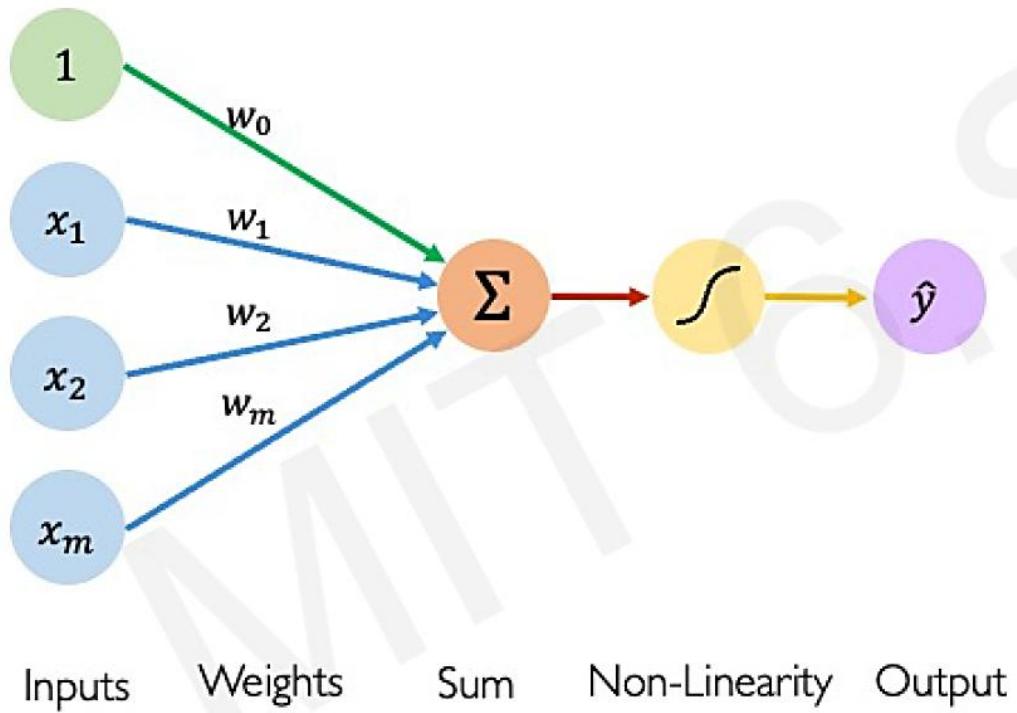
$$\hat{y} = g \left( w_0 + \sum_{i=1}^m x_i w_i \right)$$

Non-linear activation function  $\uparrow$

Bias  $\downarrow$

Output  $\downarrow$

# Perceptron

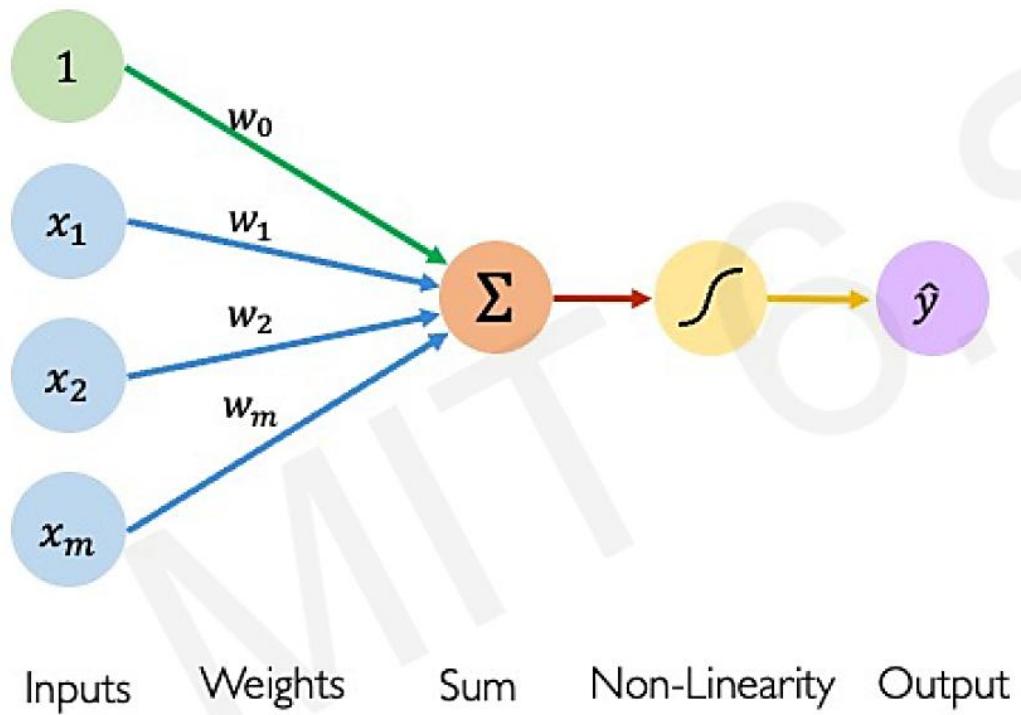


$$\hat{y} = g \left( w_0 + \sum_{i=1}^m x_i w_i \right)$$

$$\hat{y} = g ( w_0 + \mathbf{X}^T \mathbf{W} )$$

where:  $\mathbf{X} = \begin{bmatrix} x_1 \\ \vdots \\ x_m \end{bmatrix}$  and  $\mathbf{W} = \begin{bmatrix} w_1 \\ \vdots \\ w_m \end{bmatrix}$

# Perceptron

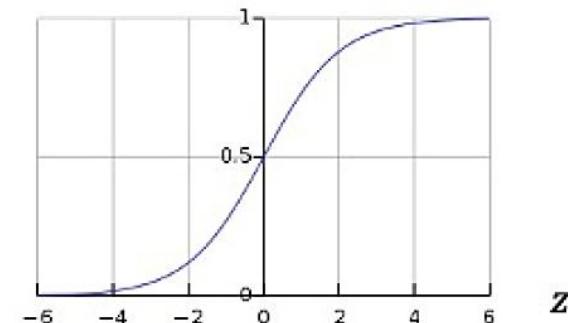


## Activation Functions

$$\hat{y} = g(w_0 + \mathbf{X}^T \mathbf{W})$$

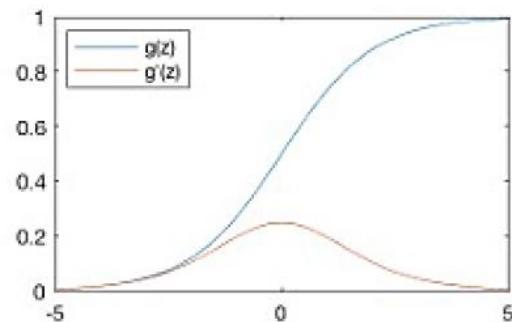
- Example: sigmoid function

$$g(z) = \sigma(z) = \frac{1}{1 + e^{-z}}$$



# Activation Functions

Sigmoid Function



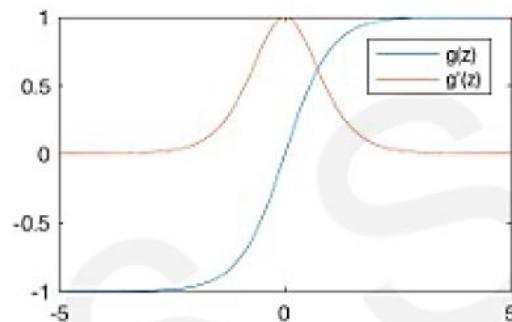
$$g(z) = \frac{1}{1 + e^{-z}}$$

$$g'(z) = g(z)(1 - g(z))$$

`tf.math.sigmoid(z)`

`torch.sigmoid(z)`

Hyperbolic Tangent



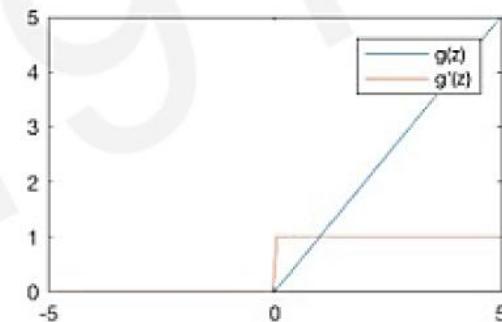
$$g(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$

$$g'(z) = 1 - g(z)^2$$

`tf.math.tanh(z)`

`torch.tanh(z)`

Rectified Linear Unit (ReLU)



$$g(z) = \max(0, z)$$

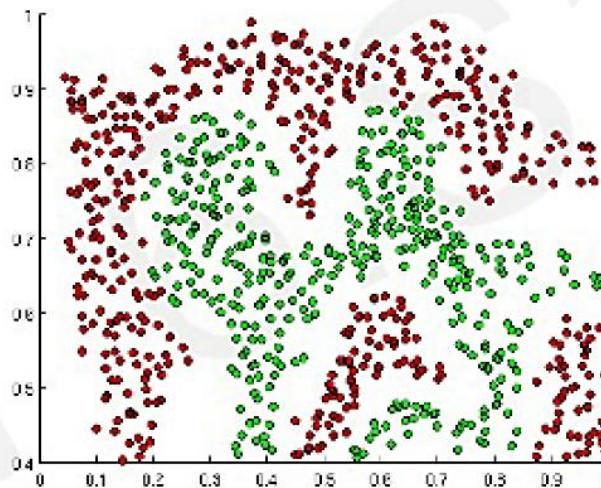
$$g'(z) = \begin{cases} 1, & z > 0 \\ 0, & \text{otherwise} \end{cases}$$

`tf.nn.relu(z)`

`torch.nn.ReLU(z)`

# Activation Functions and Non-Linearity

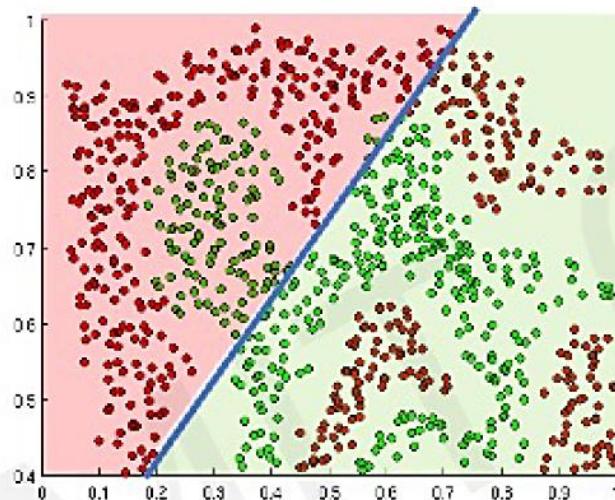
*The purpose of activation functions is to **introduce non-linearities** into the network*



What if we wanted to build a neural network to  
distinguish green vs red points?

# Activation Functions and Non-Linearity

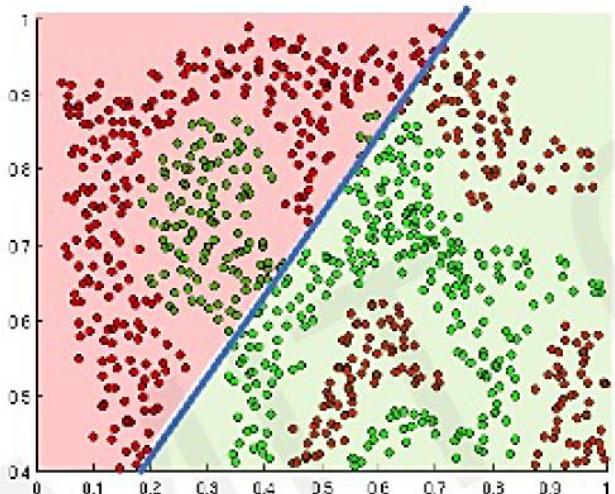
*The purpose of activation functions is to **introduce non-linearities** into the network*



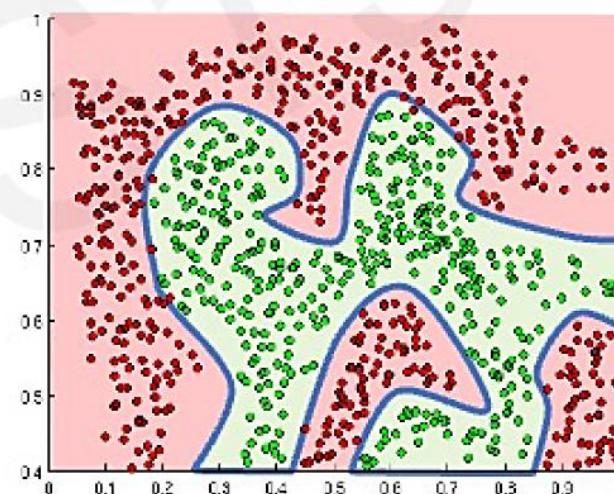
Linear activation functions produce linear decisions no matter the network size

# Activation Functions and Non-Linearity

*The purpose of activation functions is to **introduce non-linearities** into the network*

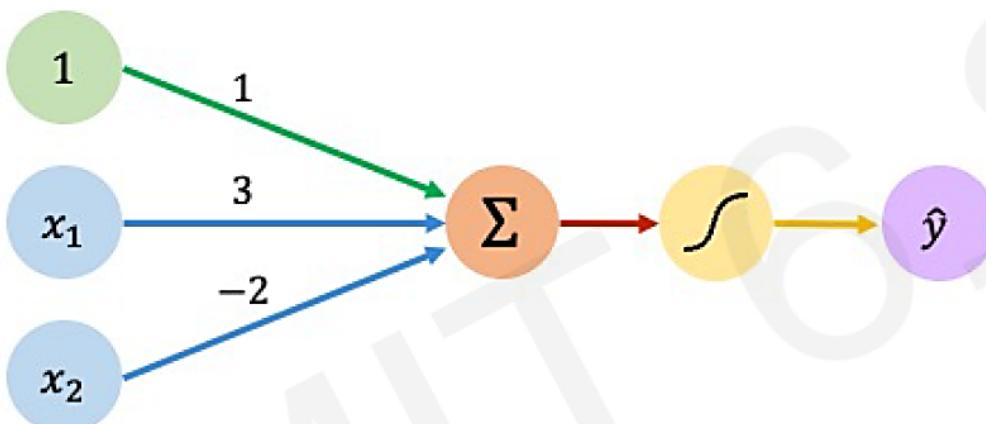


Linear activation functions produce linear decisions no matter the network size



Non-linearities allow us to approximate arbitrarily complex functions

# The Perceptron Example

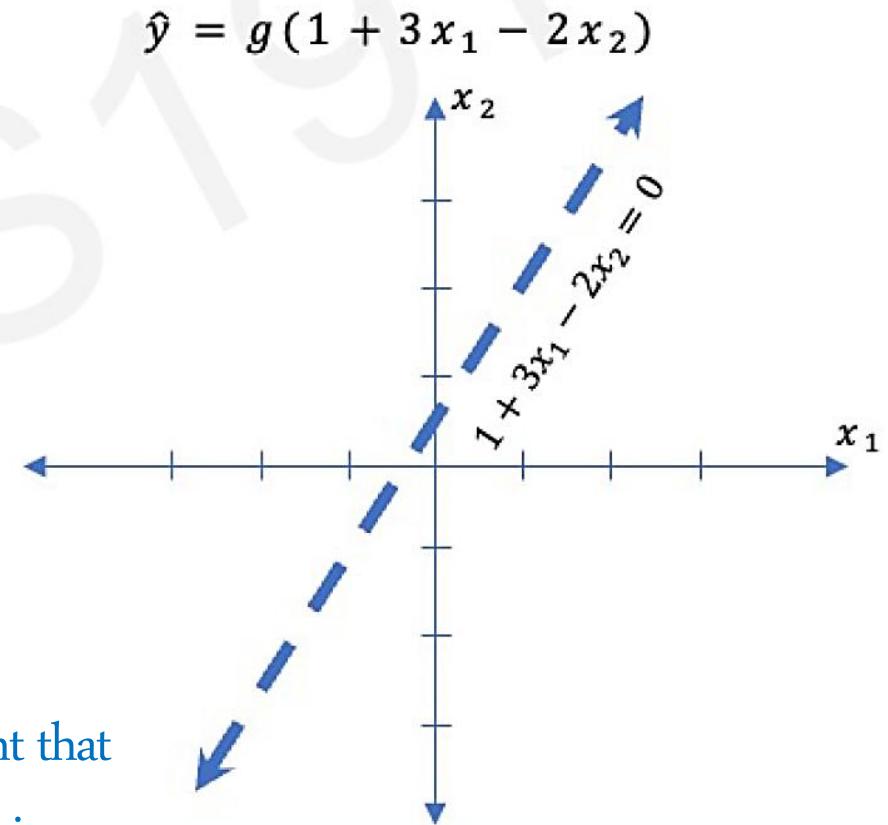
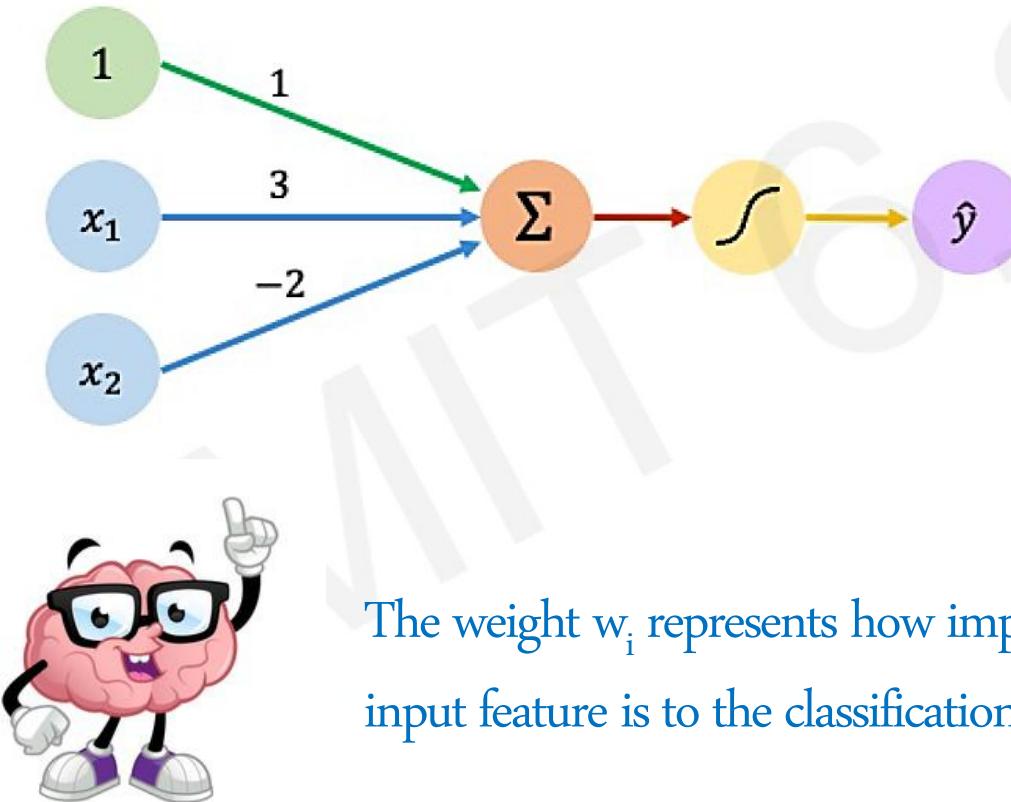


We have:  $w_0 = 1$  and  $\mathbf{w} = \begin{bmatrix} 3 \\ -2 \end{bmatrix}$

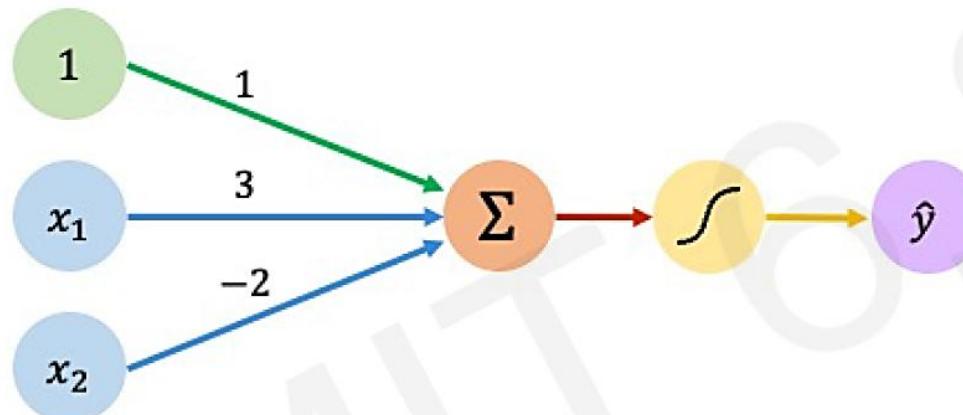
$$\begin{aligned}\hat{y} &= g(w_0 + \mathbf{x}^T \mathbf{w}) \\ &= g\left(1 + \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}^T \begin{bmatrix} 3 \\ -2 \end{bmatrix}\right) \\ \hat{y} &= g\left(1 + 3x_1 - 2x_2\right)\end{aligned}$$

This is just a line in 2D!

# The Perceptron Example

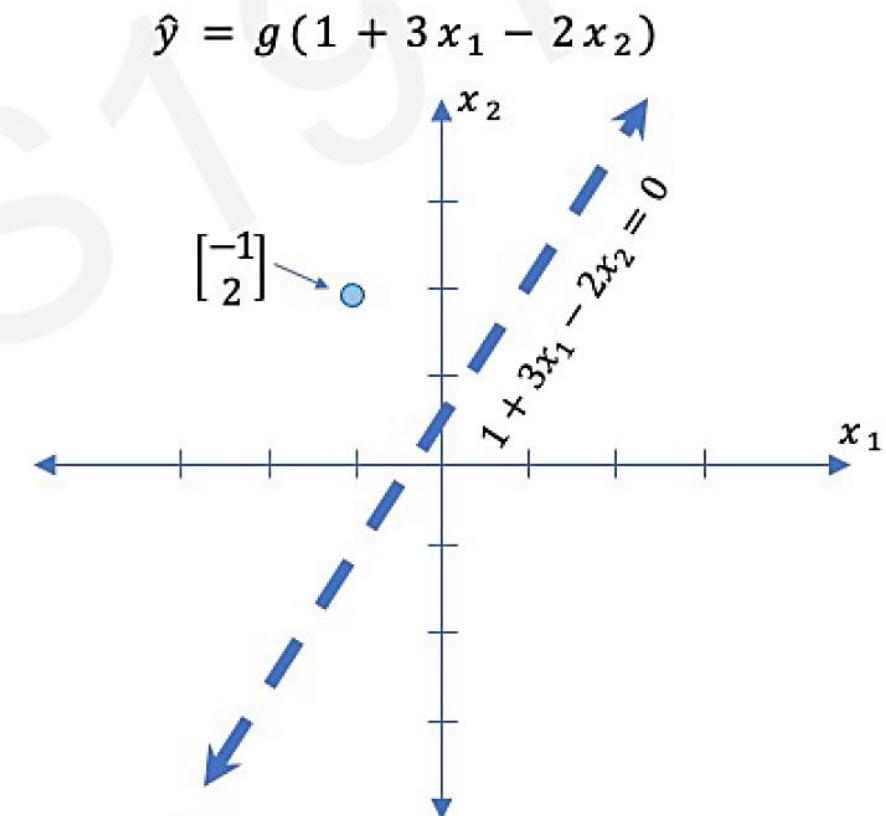


# The Perceptron Example

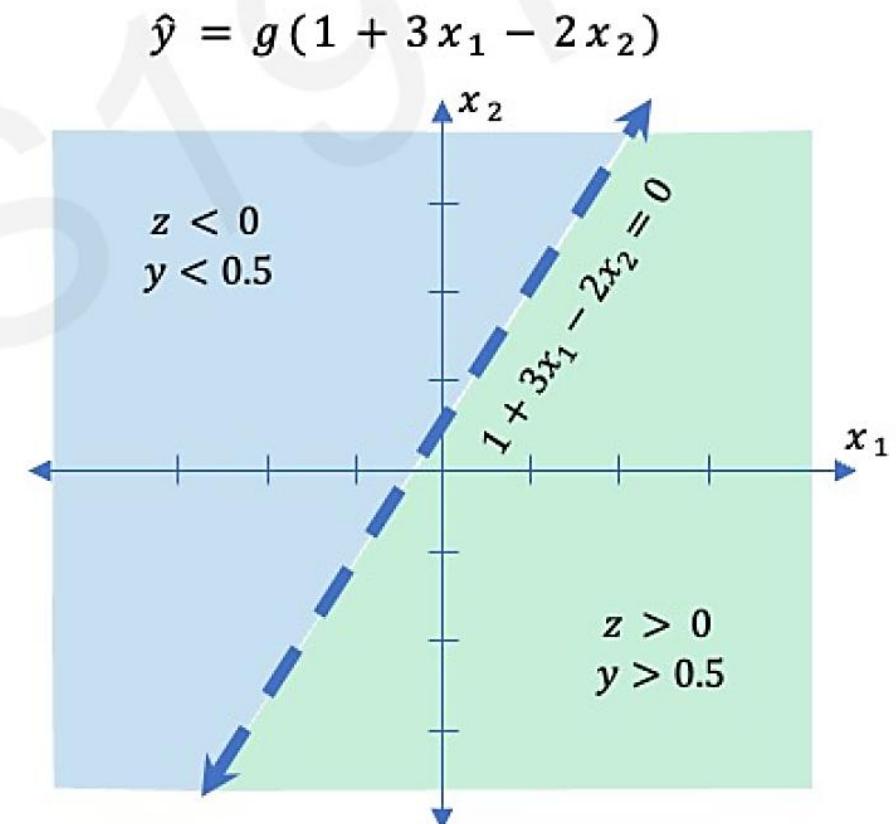
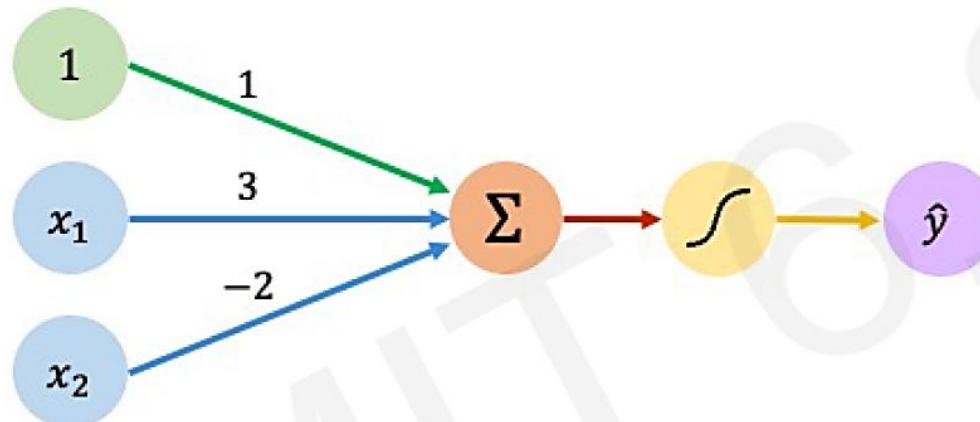


Assume we have input:  $X = \begin{bmatrix} -1 \\ 2 \end{bmatrix}$

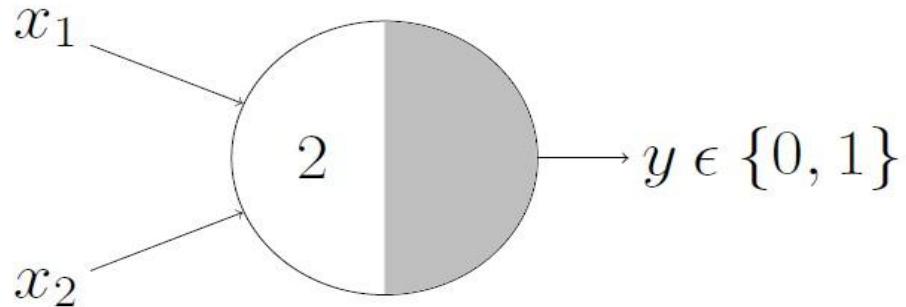
$$\begin{aligned}\hat{y} &= g(1 + (3 * -1) - (2 * 2)) \\ &= g(-6) \approx 0.002\end{aligned}$$



# The Perceptron Example

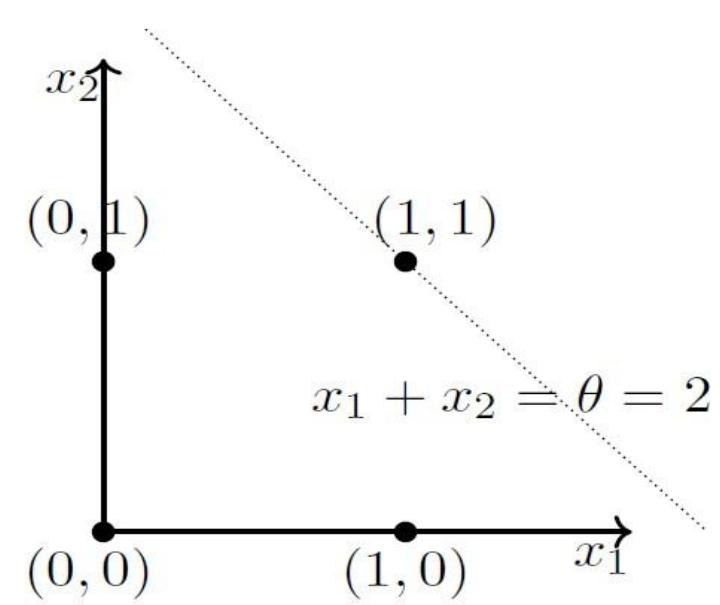
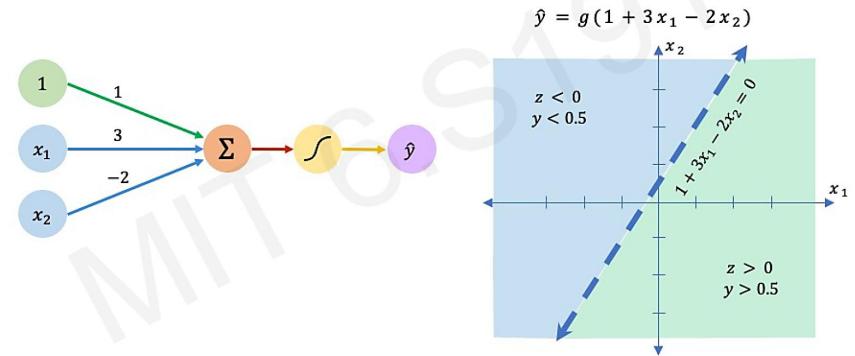


# Boolean Functions

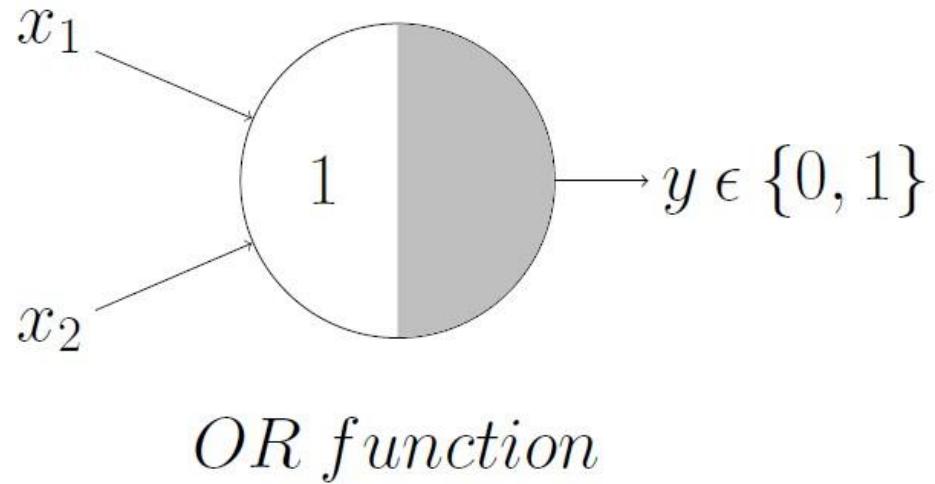


*AND function*

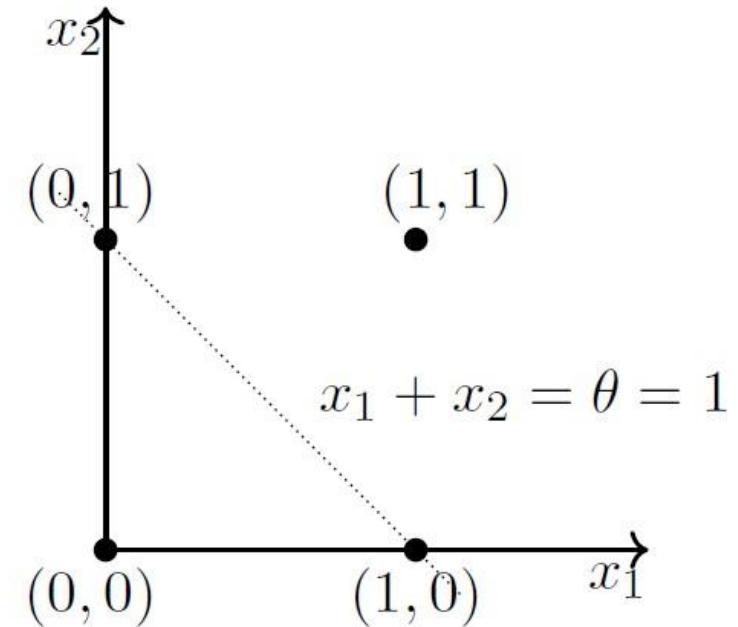
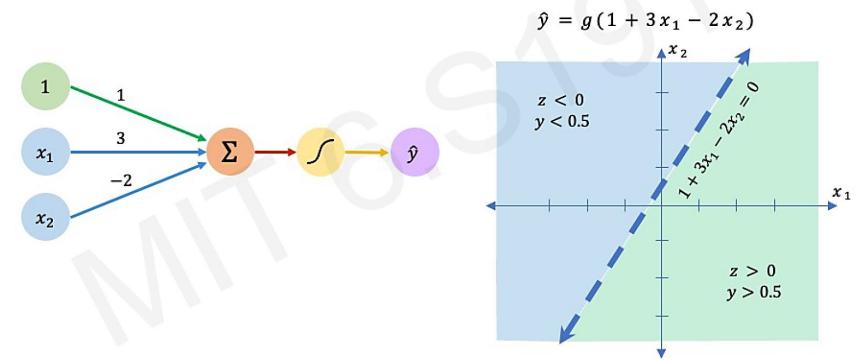
$$x_1 + x_2 = \sum_{i=1}^2 x_i \geq 2$$



# Boolean Functions



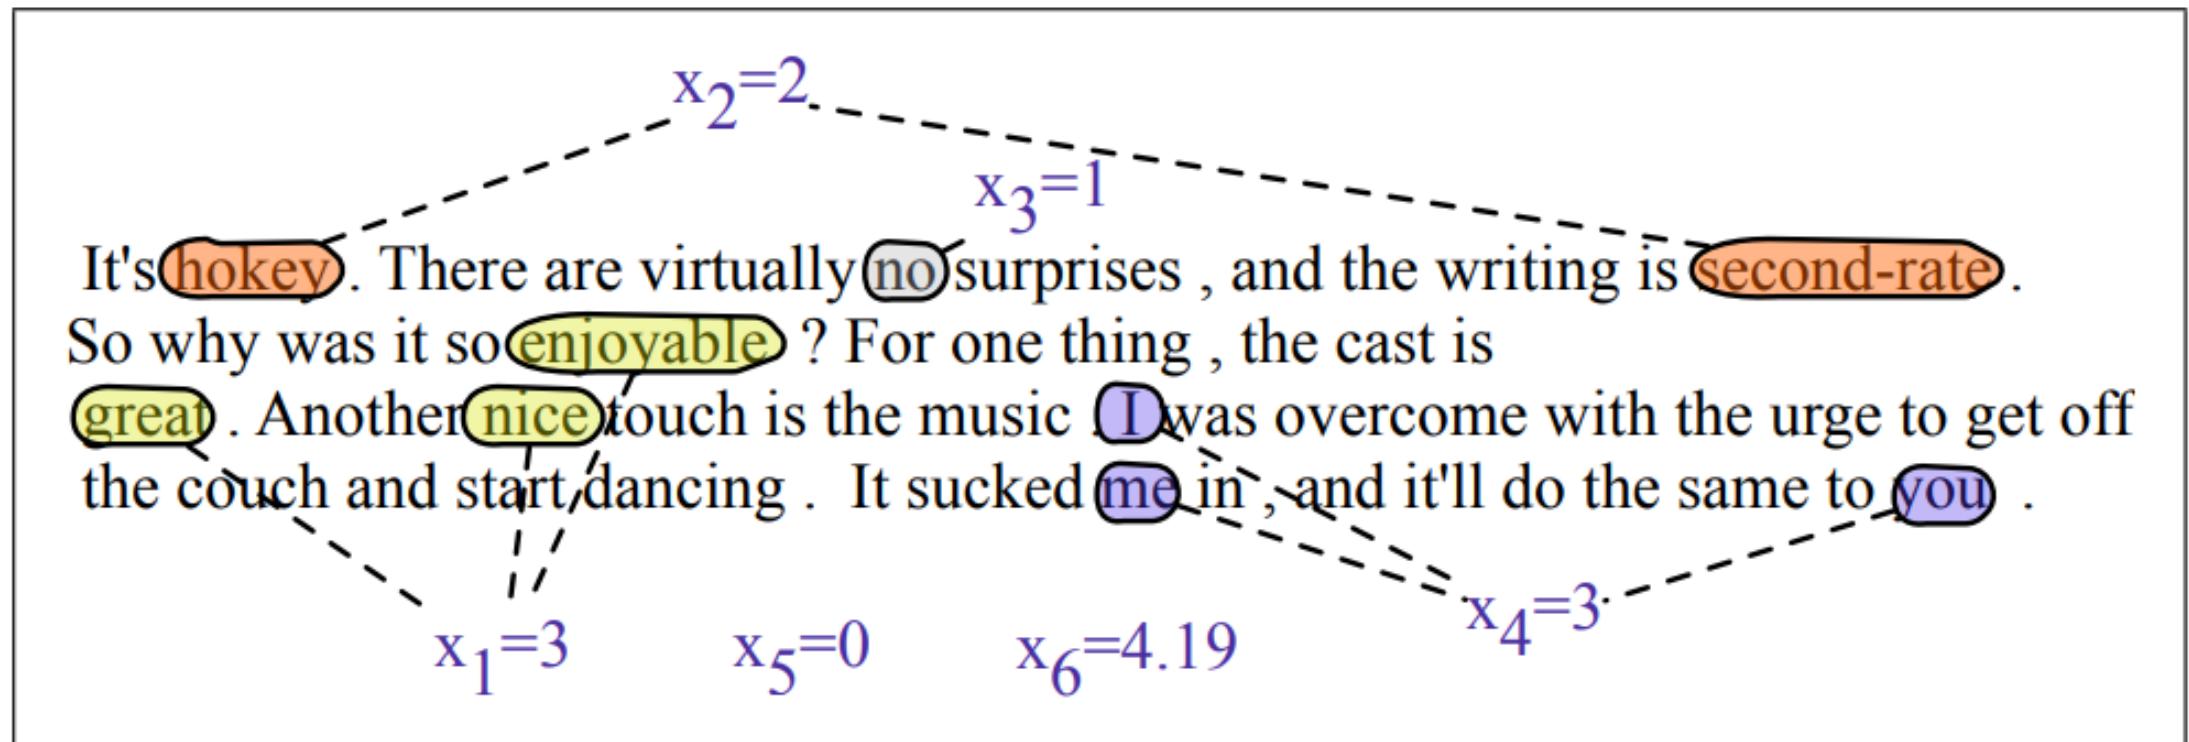
$$x_1 + x_2 = \sum_{i=1}^2 x_i \geq 1$$



# Classification with Logistic Regression

---

# A Document



**Figure 5.2** A sample mini test document showing the extracted features in the vector  $x$ .

# Features

Var	Definition	Value in Fig. 5.2
$x_1$	count(positive lexicon words $\in$ doc)	3
$x_2$	count(negative lexicon words $\in$ doc)	2
$x_3$	$\begin{cases} 1 & \text{if “no”} \in \text{doc} \\ 0 & \text{otherwise} \end{cases}$	1
$x_4$	count(1st and 2nd pronouns $\in$ doc)	3
$x_5$	$\begin{cases} 1 & \text{if “!”} \in \text{doc} \\ 0 & \text{otherwise} \end{cases}$	0
$x_6$	$\ln(\text{word count of doc})$	$\ln(66) = 4.19$

# Weights and Biases

- Let's assume the 6 weights corresponding to the 6 features are

$$w = [2.5, -5.0, -1.2, 0.5, 2.0, 0.7]$$

$$b = [0.1]$$



The weight  $w_i$  represents how important that input feature is to the classification decision

# Classification with Logistic Regression

$$p(+|x) = P(y = 1|x) = \sigma(\mathbf{w} \cdot \mathbf{x} + b)$$
$$= \sigma([2.5, -5.0, -1.2, 0.5, 2.0, 0.7] \cdot [3, 2, 1, 3, 0, 4.19] + 0.1)$$

$$= \sigma(.833)$$

$$= 0.70$$

$$p(-|x) = P(y = 0|x) = 1 - \sigma(\mathbf{w} \cdot \mathbf{x} + b)$$

$$= 0.30$$

**foreach**  $x^{(i)}$  in input  $[x^{(1)}, x^{(2)}, \dots, x^{(m)}]$

$$y^{(i)} = \sigma(\mathbf{w} \cdot \mathbf{x}^{(i)} + b)$$

$$P(y^{(1)} = 1 | x^{(1)}) = \sigma(\mathbf{w} \cdot \mathbf{x}^{(1)} + b)$$

$$P(y^{(2)} = 1 | x^{(2)}) = \sigma(\mathbf{w} \cdot \mathbf{x}^{(2)} + b)$$

$$P(y^{(3)} = 1 | x^{(3)}) = \sigma(\mathbf{w} \cdot \mathbf{x}^{(3)} + b)$$

$$\mathbf{X} = \begin{bmatrix} x_1^{(1)} & x_2^{(1)} & \dots & x_f^{(1)} \\ x_1^{(2)} & x_2^{(2)} & \dots & x_f^{(2)} \\ x_1^{(3)} & x_2^{(3)} & \dots & x_f^{(3)} \\ \dots & & & \end{bmatrix}$$

$$\hat{y}^{(1)} = [x_1^{(1)}, x_2^{(1)}, \dots, x_f^{(1)}] \cdot [w_1, w_2, \dots, w_f] + b$$

$$\begin{array}{ccccccccc} \mathbf{y} & = & \mathbf{X} & & \mathbf{w} & + & \mathbf{b} \\ & & (m \times 1) & & (m \times f) & (f \times 1) & (m \times 1) \end{array}$$

## Processing many examples at once

# Cost Function

---

# Cost Function

Training set:  $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})\}$

m examples

$$x \in \begin{bmatrix} x_0 \\ x_1 \\ \dots \\ x_n \end{bmatrix} \quad x_0 = 1, y \in \{0, 1\}$$

$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

How to choose parameters  $\theta$  ?

# Cost Function

$L(\hat{y}, y)$  = How much  $\hat{y}$  differs from the true  $y$ .



# A good cost function?

1. Punish incorrect answers with high cost

$$y=1, h\theta(x) \rightarrow 0$$

$$y=0, h\theta(x) \rightarrow 1$$

2. Reward correct answers with a low cost

$$y=1, h\theta(x) \rightarrow 1$$

$$y=0, h\theta(x) \rightarrow 0$$



# The cross-entropy loss function

$$\begin{aligned} J(\theta) &= \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_\theta(x^{(i)}), y^{(i)}) \\ &= -\frac{1}{m} \left[ \sum_{i=1}^m y^{(i)} \log h_\theta(x^{(i)}) + (1 - y^{(i)}) \log (1 - h_\theta(x^{(i)})) \right] \end{aligned}$$

$L(\hat{y}, y)$  = How much  $\hat{y}$  differs from the true  $y$ .



- Rewards  $y = 1, h_\theta(x) = 1$ , with cost = 0
- Rewards  $y = 1, h_\theta(x) \rightarrow 0$  even further with asymptotically smaller costs

**y=1**

cost =  $\log(h_\theta(x))$

- Rewards  $y = 1, h_\theta(x) = 1$ , with cost = 0
- Punishes  $y = 1, h_\theta(x) \rightarrow 0$  with asymptotically higher costs

**y=1**

cost =  $-\log(h_\theta(x))$

- Rewards  $y = 0, h_\theta(x) = 0$ , with cost = 0
- Rewards  $y = 0, h_\theta(x) \rightarrow 1$  even further with asymptotically smaller costs

**y=0**

$$\text{cost} = \log(1 - h_\theta(x))$$

- Rewards  $y = 0, h_\theta(x) = 0$ , with cost = 0
- Punishes  $y = 0, h_\theta(x) \rightarrow 1$  with asymptotically higher costs

**y=0**

$$\text{cost} = -\log(1 - h_\theta(x))$$

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_\theta(x^{(i)}), y^{(i)})$$

$$\text{Cost}(h_\theta(x), y) = \begin{cases} -\log(h_\theta(x)) & \text{if } y = 1 \\ -\log(1 - h_\theta(x)) & \text{if } y = 0 \end{cases}$$

$$cost(h_{\theta}(x), y) = -y \log(h_{\theta}(x)) - (1 - y) \log(1 - h_{\theta}(x))$$

$$L_{CE}(\hat{y}, y) = -[y \log \sigma(\mathbf{w} \cdot \mathbf{x} + b) + (1 - y) \log (1 - \sigma(\mathbf{w} \cdot \mathbf{x} + b))]$$

$$\begin{aligned} J(\theta) &= \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^{(i)}), y^{(i)}) \\ &= -\frac{1}{m} \left[ \sum_{i=1}^m y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log (1 - h_{\theta}(x^{(i)})) \right] \end{aligned}$$

To fit parameters  $\theta$ :

$$\min_{\theta} J(\theta)$$

To make a prediction given new  $x$ :

$$\text{Output } h_{\theta}(x) = \frac{1}{1+e^{-\theta^T x}} \quad \text{gives } p(y = 1 | x; \theta)$$

$$\begin{aligned}L_{\text{CE}}(\hat{y}, y) &= -[y \log \sigma(\mathbf{w} \cdot \mathbf{x} + b) + (1 - y) \log (1 - \sigma(\mathbf{w} \cdot \mathbf{x} + b))] \\&= -[\log \sigma(\mathbf{w} \cdot \mathbf{x} + b)] \\&= -\log(.70) \\&= .36\end{aligned}$$

$$\begin{aligned}L_{\text{CE}}(\hat{y}, y) &= -[y \log \sigma(\mathbf{w} \cdot \mathbf{x} + b) + (1 - y) \log (1 - \sigma(\mathbf{w} \cdot \mathbf{x} + b))] \\&= -[\log (1 - \sigma(\mathbf{w} \cdot \mathbf{x} + b))] \\&= -\log (.30) \\&= 1.2\end{aligned}$$

# Gradient Descent

---

# Gradient Descent

The goal is to find the set of weights which minimizes the loss function, averaged over all examples:

$$\hat{\theta} = \operatorname{argmin}_{\theta} \frac{1}{m} \sum_{i=1}^m L_{\text{CE}}(f(x^{(i)}; \theta), y^{(i)})$$

Want  $\min_{\theta} J(\theta)$ :

Repeat {

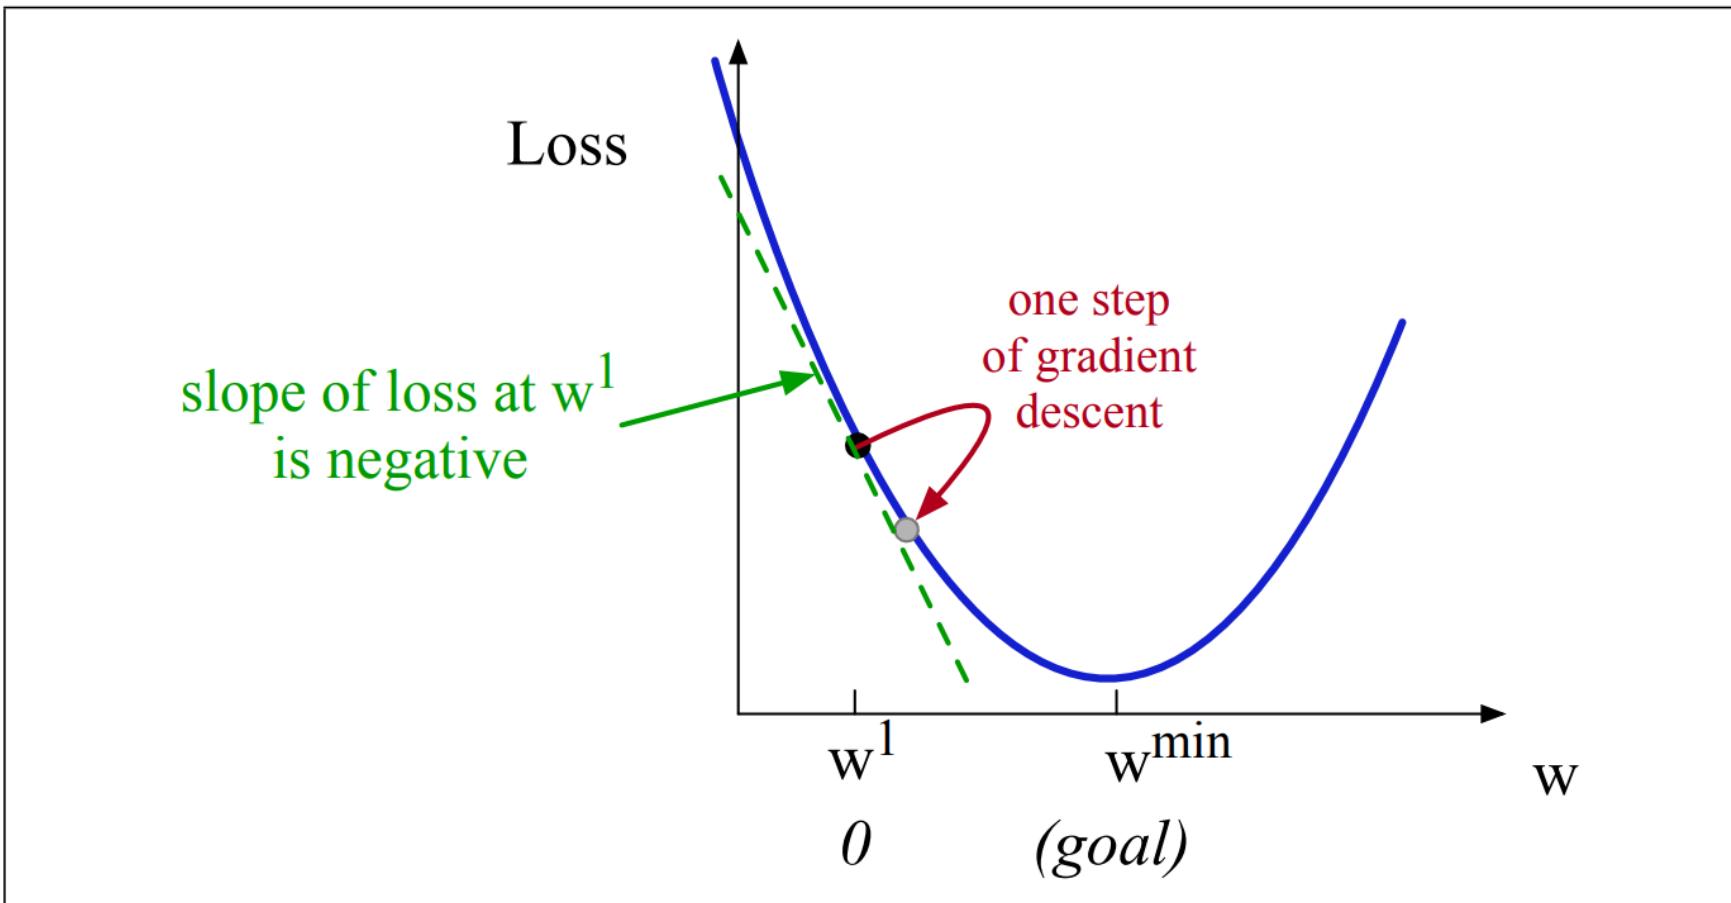
$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

}

$$w^{t+1} = w^t - \eta \frac{d}{dw} L(f(x; w), y)$$

(simultaneously update all  $\theta_j$ )

# Gradient Descent



---

```
function STOCHASTIC GRADIENT DESCENT( $L()$ ,  $f()$ ,  $x$ ,  $y$ ) returns  $\theta$ 
    # where: L is the loss function
    #      f is a function parameterized by  $\theta$ 
    #      x is the set of training inputs  $x^{(1)}$ ,  $x^{(2)}$ , ...,  $x^{(m)}$ 
    #      y is the set of training outputs (labels)  $y^{(1)}$ ,  $y^{(2)}$ , ...,  $y^{(m)}$ 

     $\theta \leftarrow 0$       # (or small random values)
    repeat til done  # see caption
        For each training tuple  $(x^{(i)}, y^{(i)})$  (in random order)
            1. Optional (for reporting):      # How are we doing on this tuple?
                Compute  $\hat{y}^{(i)} = f(x^{(i)}; \theta)$  # What is our estimated output  $\hat{y}$ ?
                Compute the loss  $L(\hat{y}^{(i)}, y^{(i)})$  # How far off is  $\hat{y}^{(i)}$  from the true output  $y^{(i)}$ ?
            2.  $g \leftarrow \nabla_{\theta} L(f(x^{(i)}; \theta), y^{(i)})$       # How should we move  $\theta$  to maximize loss?
            3.  $\theta \leftarrow \theta - \eta g$                       # Go the other way instead
    return  $\theta$ 
```

---

# Batch Gradient Descent

$$J(\theta) = -\frac{1}{m} \left[ \sum_{i=1}^m y^{(i)} \log h_\theta(x^{(i)}) + (1 - y^{(i)}) \log (1 - h_\theta(x^{(i)})) \right]$$

Want  $\min_{\theta} J(\theta)$ :

Repeat {

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

}

(simultaneously update all  $\theta_j$ )

(see SLP3 for derivation)

$$\begin{aligned} \frac{\partial L_{\text{CE}}(\hat{y}, y)}{\partial w_j} &= [\sigma(\mathbf{w} \cdot \mathbf{x} + b) - y] x_j \\ &= (\hat{y} - y) x_j \end{aligned}$$

$$\frac{\partial}{\partial \theta_j} J(\theta) = \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) x^{(i)}$$



# Batch Gradient Descent

$$J(\theta) = -\frac{1}{m} \left[ \sum_{i=1}^m y^{(i)} \log h_\theta(x^{(i)}) + (1 - y^{(i)}) \log (1 - h_\theta(x^{(i)})) \right]$$

Want  $\min_{\theta} J(\theta)$ :

Repeat {

$$\theta_j := \theta_j - \alpha \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

(simultaneously update all  $\theta_j$ )

}

Algorithm looks identical to linear regression!

$$\nabla L(f(x; \theta), y) = \begin{bmatrix} \frac{\partial}{\partial w_1} L(f(x; \theta), y) \\ \frac{\partial}{\partial w_2} L(f(x; \theta), y) \\ \vdots \\ \frac{\partial}{\partial w_n} L(f(x; \theta), y) \\ \frac{\partial}{\partial b} L(f(x; \theta), y) \end{bmatrix}$$



# Gradient Descent Example

$x_1 = 3$  (count of positive lexicon words)

$x_2 = 2$  (count of negative lexicon words)

$w_1 = w_2 = b = 0$

$\eta = 0.1$

$$\theta^{t+1} = \theta^t - \eta \nabla_{\theta} L(f(x^{(i)}; \theta), y^{(i)})$$

$$\nabla_{w,b} L = \begin{bmatrix} \frac{\partial L_{\text{CE}}(\hat{y}, y)}{\partial w_1} \\ \frac{\partial L_{\text{CE}}(\hat{y}, y)}{\partial w_2} \\ \frac{\partial L_{\text{CE}}(\hat{y}, y)}{\partial b} \end{bmatrix} = \begin{bmatrix} (\sigma(\mathbf{w} \cdot \mathbf{x} + b) - y)x_1 \\ (\sigma(\mathbf{w} \cdot \mathbf{x} + b) - y)x_2 \\ \sigma(\mathbf{w} \cdot \mathbf{x} + b) - y \end{bmatrix} = \begin{bmatrix} (\sigma(0) - 1)x_1 \\ (\sigma(0) - 1)x_2 \\ \sigma(0) - 1 \end{bmatrix} = \begin{bmatrix} -0.5x_1 \\ -0.5x_2 \\ -0.5 \end{bmatrix} = \begin{bmatrix} -1.5 \\ -1.0 \\ -0.5 \end{bmatrix}$$

$$\theta^1 = \begin{bmatrix} w_1 \\ w_2 \\ b \end{bmatrix} - \eta \begin{bmatrix} -1.5 \\ -1.0 \\ -0.5 \end{bmatrix} = \begin{bmatrix} .15 \\ .1 \\ .05 \end{bmatrix}$$

So, after one step of gradient descent, the weights have shifted to be:  $w_1 = .15$ ,  $w_2 = .1$ , and  $b = .05$



# Softmax

---

# Softmax

- For a vector  $z$  of dimensionality  $K$ , the softmax is defined as:

$$\text{softmax}(z_i) = \frac{\exp(z_i)}{\sum_{j=1}^K \exp(z_j)} \quad 1 \leq i \leq K$$

- The softmax of an input vector  $z = [z_1, z_2, \dots, z_K]$  is thus a vector itself:

$$\text{softmax}(z) = \left[ \frac{\exp(z_1)}{\sum_{i=1}^K \exp(z_i)}, \frac{\exp(z_2)}{\sum_{i=1}^K \exp(z_i)}, \dots, \frac{\exp(z_K)}{\sum_{i=1}^K \exp(z_i)} \right]$$

# Softmax

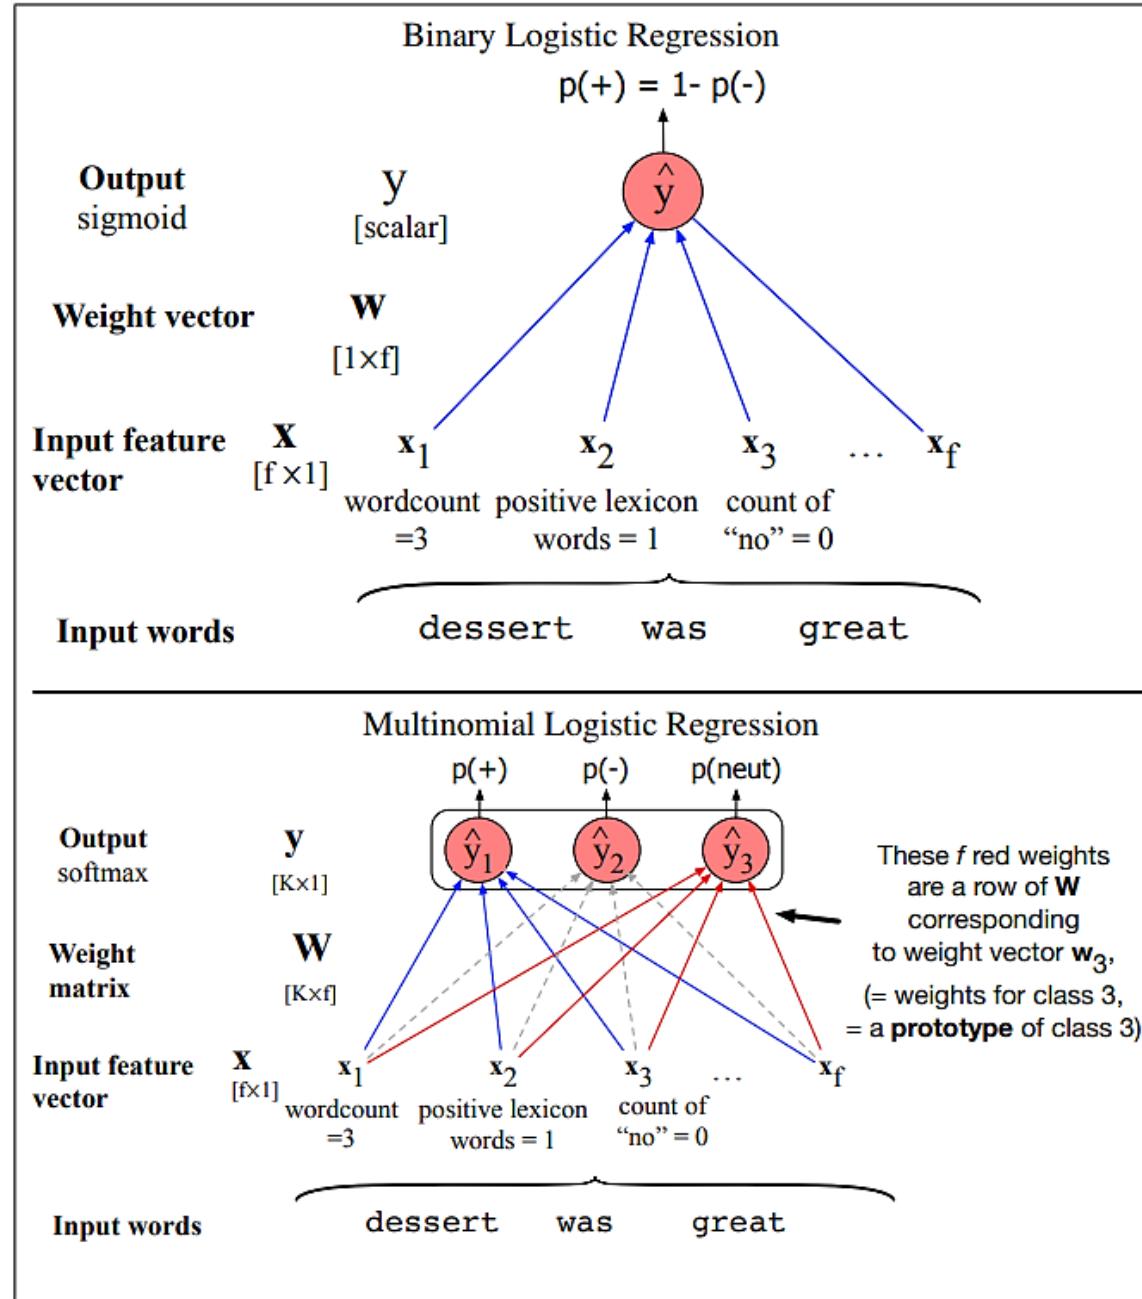
$$\text{softmax}(\mathbf{z}) = \left[ \frac{\exp(z_1)}{\sum_{i=1}^K \exp(z_i)}, \frac{\exp(z_2)}{\sum_{i=1}^K \exp(z_i)}, \dots, \frac{\exp(z_K)}{\sum_{i=1}^K \exp(z_i)} \right]$$

$$\mathbf{z} = [0.6, 1.1, -1.5, 1.2, 3.2, -1.1]$$

The resulting (rounded) softmax( $\mathbf{z}$ ) is:

$$[0.05, 0.09, 0.01, 0.1, 0.74, 0.01]$$

A neural network can  
be viewed as a series of  
logistic regression  
classifiers stacked on  
top of each other



# Sources

- <https://web.stanford.edu/~jurafsky/slp3/2.pdf>
- <https://web.stanford.edu/~jurafsky/slp3/3.pdf>
- **Machine Learning for Intelligent Systems**, Kilian Weinberger, Cornell, Lectures 3-6,  
[https://www.cs.cornell.edu/courses/cs4780/2018fa/lectures/lecture note03.html](https://www.cs.cornell.edu/courses/cs4780/2018fa/lectures/lecture_note03.html)
- **Prof. Mitesh M. Khapra** (<https://www.cse.iitm.ac.in/~miteshk/>) on NPTEL's (<http://nptel.ac.in/>) Deep Learning course ([https://onlinecourses.nptel.ac.in/noc18 cs41/preview](https://onlinecourses.nptel.ac.in/noc18_cs41/preview))
- **Perceptrons. An Introduction to Computational Geometry.** Marvin Minsky and Seymour Papert. M.I.T. Press, Cambridge, Mass., 1969. <https://science.sciencemag.org/content/165/3895/780>