



## Module: Data Exploration in Azure Databricks and Introduction to Delta Lake

Microsoft Services



## Module Overview

- Lesson 1: Data Exploration with Azure Databricks
- Lesson 2: Introduction to Delta Lake

## Lesson 1: Data Exploration with Azure Databricks

After completing this lesson, you will be able to:

- Understand Apache Spark DataFrame and Dataset
- Understand data exploration techniques using Python or Scala

# Spark Interfaces

## Resilient Distributed Dataset (RDD)

Spark RDD is a resilient, partitioned, distributed and immutable collection of data.

## DataFrame

Distributed collection of data organized into named columns. It is conceptually equivalent to a table in a relational database or a data frame in R/Python, but with richer optimizations under the hood

## Dataset

A Dataset is a strongly-typed, immutable collection of objects that are mapped to a relational schema.

An extension of the DataFrame API that provides a *type-safe, object-oriented programming interface*.

All that you are going to do in Apache Spark is to read some data from a source and load it into Spark. You will then process the data and hold the intermediate results, and finally write the results back to a destination. But in this process, you need a data structure to hold the data in Spark. We have three alternatives to hold data in Spark.

RDD

Data Frame

Dataset

Apache Spark 2.x recommends to use the last two and avoid using RDDs. However, there is a critical fact to note about RDDs. Data Frames and Datasets, both of them are ultimately compiled down to an RDD. So, under the hood, everything in Spark is an RDD. And for that reason, I will start with RDDs and try to explain the mechanics of parallel processing.

## Spark RDD

Let's define the RDD. The name stands for Resilient Dristributed Dataset. However, I can describe a RDD as below.

Spark RDD is a resilient, partitioned, distributed and immutable collection of data. Let's

quickly review this description.

Collection of data - RDDs hold data and appears to be a Scala Collection.

Resilient - RDDs can recover from a failure, so they are fault tolerant.

Partitioned - Spark breaks the RDD into smaller chunks of data. These pieces are called partitions.

Distributed - Instead of keeping those partitions on a single machine, Spark spreads them across the cluster. So they are a distributed collection of data.

Immutable - Once defined, you can't change a RDD. So Spark RDD is a read-only data structure.

### When to use RDDs?

Consider these scenarios or common use cases for using RDDs when:

you want low-level transformation and actions and control on your dataset;

your data is unstructured, such as media streams or streams of text;

you want to manipulate your data with functional programming constructs than domain specific expressions;

you don't care about imposing a schema, such as columnar format, while processing or accessing data attributes by name or column; and

you can forgo some optimization and performance benefits available with

DataFrames and Datasets for structured and semi-structured data.

More Information: <https://databricks.com/blog/2016/07/14/a-tale-of-three-apache-spark-apis-rdds-dataframes-and-datasets.html>

### What Are DataFrames?

In Spark, a DataFrame is a distributed collection of data organized into named columns. It is conceptually equivalent to a table in a relational database or a data frame in R/Python, but with richer optimizations under the hood. DataFrames can be constructed from a wide array of sources such as: structured data files, tables in Hive, external databases, or existing RDDs.

Similar to RDDs, DataFrames are evaluated lazily. That is to say, computation only happens when an action (e.g. display result, save output) is required. This allows their executions to be optimized, by applying techniques such as predicate push-downs and bytecode generation, as explained later in the section "Under the Hood: Intelligent Optimization and Code Generation". All DataFrame operations are also automatically parallelized and distributed on clusters.

More Information: <https://databricks.com/blog/2015/02/17/introducing-dataframes-in-spark-for-large-scale-data-science.html>

## DataSets

A Dataset is a strongly-typed, immutable collection of objects that are mapped to a relational schema. At the core of the Dataset API is a new concept called an encoder, which is responsible for converting between JVM objects and tabular representation. The tabular representation is stored using Spark's internal Tungsten binary format, allowing for operations on serialized data and improved memory utilization. Spark 1.6 comes with support for automatically generating encoders for a wide variety of types, including primitive types (e.g. String, Integer, Long), Scala case classes, and Java Beans.

More Information: <https://databricks.com/blog/2016/01/04/introducing-apache-spark-datasets.html>

In many scenarios, especially with the performance optimizations embedded in DataFrames and Datasets, it will not be necessary to work with RDDs. But it is important to understand the RDD abstraction because:

The RDD is the underlying infrastructure that allows Spark to run so fast and provide data lineage.

If you are diving into more advanced components of Spark, it may be necessary to use RDDs.

The visualizations within the [Spark UI](#) reference RDDs.

## DataFrames

- DataFrame is a distributed collection of data organized into named columns.
- Conceptually equivalent to a table in a relational database or a data frame in R/Python, but with richer optimizations.
- DataFrames can be constructed from a wide array of sources such as: structured data files, tables in Hive, external databases, or existing RDDs.
- DataFrames are evaluated lazily, meaning, computation only happens when an action (e.g. display result, save output) is required.

## Loading Data in DataFrames

```
%python
# Use the Spark CSV datasource with options specifying:
# - First line of file is a header
# - Automatically infer the schema of the data

data = spark.read.format("csv") \
    .option("header", "true") \
    .option("inferSchema", "true") \
    .load("/databricks-datasets/samples/population-vs-\
price/data_geo.csv")

... do something
```

The easiest way to start working with DataFrames is to use an example Azure Databricks dataset available in the /databricks-datasets folder accessible within the Azure Databricks workspace. To access the file that compares city population versus median sale prices of homes, load the file /databricks-datasets/samples/population-vs-price/data\_geo.csv.

# Viewing DataFrames

Using Spark Command `take()` to view raw records

```
%python data.take(10) #view 10 records of DataFrame
```

▶ (1) Spark Jobs

`Out[3]:`

```
[Row(2014 rank=101, City=u'Birmingham', State=u'Alabama', State Code=u'AL', 2014 Population estimate=212247, 2015 median sales price=162.9),
 Row(2014 rank=125, City=u'Huntsville', State=u'Alabama', State Code=u'AL', 2014 Population estimate=188226, 2015 median sales price=157.7),
 Row(2014 rank=122, City=u'Mobile', State=u'Alabama', State Code=u'AL', 2014 Population estimate=194675, 2015 median sales price=122.5),
```

Using `display()` to view in tabular mode

```
%python display(data)
```

▶ (2) Spark Jobs

2014 rank	City	State	State Code	2014 Population estimate	2015 median sales price
101	Birmingham	Alabama	AL	212247	162.9
125	Huntsville	Alabama	AL	188226	157.7
122	Mobile	Alabama	AL	194675	122.5

Now that you have created the data DataFrame, you can quickly access the data using standard Spark commands such as `take()`. For example, you can use the command `data.take(10)` to view the first ten rows of the data DataFrame. Because this is a SQL notebook, the next few commands use the [%python magic command](#).

To view this data in a tabular format, you can use the Azure Databricks `display()` command instead of exporting the data to a third-party tool.

## Run SQL Queries

Before you can issue SQL queries, you must save your data DataFrame as a temporary table:

```
%python  
# Register table so it is accessible via SQL Context  
data.createOrReplaceTempView("data_geo")
```

Then, in a new cell, specify a SQL query to list the 2015 median sales price by state:

```
select 'State Code', '2015 median sales price' from data_geo
```

Using the dataframe "data" we will create a temporary view/table named "data\_geo". Once we have a tempview created, we can run sql queries against "data\_geo".

Temporary tables persist for the duration of the spark session.

## Datasets

- The Apache Spark Dataset API provides a type-safe, object-oriented programming interface
- DataFrame is an alias for an untyped Dataset [Row]
- Datasets provide compile-time type safety
- The Dataset API also offers high-level domain-specific language operations
- Datasets are only available in Scala/Java

The Apache Spark [Dataset API](#) provides a type-safe, object-oriented programming interface. DataFrame is an alias for an untyped Dataset [Row]. Datasets provide compile-time type safety—which means that production applications can be checked for errors before they are run—and they allow direct operations over user-defined classes. The Dataset API also offers high-level domain-specific language operations like sum(), avg(), join(), select(), groupBy(), making the code a lot easier to express, read, and write.

Must Read to understand DataSet:

<https://spark.apache.org/docs/latest/api/java/org/apache/spark/sql/Dataset.html>

## Datasets

- Operations available on Datasets are divided into transformations and actions.
  - **Transformations** are the ones that produce new Datasets
    - Example transformations include map, filter, select, and aggregate (groupBy).
  - **Actions** are the ones that trigger computation and return results.
    - Example actions count, show, or writing data out to file systems.
- Datasets are "lazy", i.e. computations are only triggered when an action is invoked.

The Apache Spark [Dataset API](#) provides a type-safe, object-oriented programming interface. DataFrame is an alias for an untyped Dataset [Row]. Datasets provide compile-time type safety—which means that production applications can be checked for errors before they are run—and they allow direct operations over user-defined classes. The Dataset API also offers high-level domain-specific language operations like sum(), avg(), join(), select(), groupBy(), making the code a lot easier to express, read, and write.

Must Read to understand DataSet:

<https://spark.apache.org/docs/latest/api/java/org/apache/spark/sql/Dataset.html>

## Datasets (cont...)

- To select a column from the Dataset:

```
val ageCol = people("age") // in Scala  
Column ageCol = people.col("age"); // in Java
```

- . Transformation Example(Filtering)

```
people.filter("age > 30")  
  .join(department, people("deptId") === department("id"))  
  .groupBy(department("name"), people("gender"))  
  .agg(avg(people("salary")), max(people("age")))
```

<https://spark.apache.org/docs/latest/api/java/org/apache/spark/sql/Dataset.html>

## Datasets (cont...)

```
case class Person(name: String, age: Long)

// Encoders are created for case classes
val caseClassDS = Seq(Person("Andy", 32)).toDS()
caseClassDS.show()

val path = "examples/src/main/resources/people.json"
val peopleDS = spark.read.json(path).as[Person]
peopleDS.show()
```

<https://spark.apache.org/docs/latest/sql-getting-started.html#creating-datasets>

## Datasets (cont...)

```
// Register the DataFrame as a SQL temporary view
df.createOrReplaceTempView("people")

val sqlDF = spark.sql("SELECT * FROM people")
sqlDF.show()
// +---+-----+
// | age| name|
// +---+-----+
// |null|Michael|
// | 30| Andy|
// | 19| Justin|
// +---+-----+
```

<https://spark.apache.org/docs/latest/sql-getting-started.html#running-sql-queries-programmatically>

## Datasets (cont...)

- To efficiently support domain-specific objects, an Encoder is required.
  - The encoder maps the domain specific type T to Spark's internal type system.
  - For example, given a class Person with two fields, name (string) and age (int), an encoder is used to tell Spark to generate code at runtime to serialize the Person object into a binary structure.
- This binary structure often has much lower memory footprint as well as are optimized for efficiency in data processing (e.g. in a columnar format).

# Lab: Data Exploration with Azure Databricks

M03 L01 Lab 01  
30 mins



Refer to M03\\_\\_L01\\_\\_Lab01.ipynb file for Databricks notebook code. Remember to import the notebook to your workspace and read the directions before executing the code

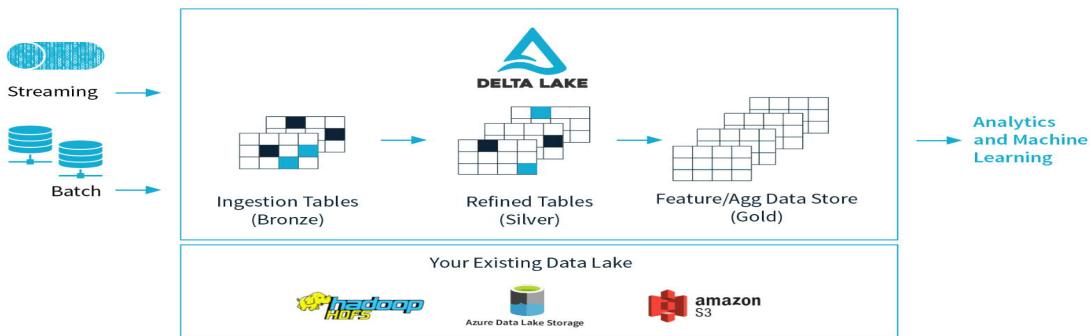
## Lesson 2: Introduction to Delta Lake

After completing this learning unit, you will be able to:

- Explore the main features of Delta Lake
- Learn how to read from and write to Delta tables from interactive queries.
- Learn how to optimize Delta Lake performances with file management

## Databricks Delta Lake Review

- Delta Lake is a storage layer that brings scalable, ACID transactions to [Apache Spark](#) and other big-data engines.
- Delta Lake runs on top of your existing cloud storage and is fully compatible with Apache Spark APIs



[Delta Lake](#) is an [open source storage layer](#) that brings reliability to data lakes. Delta Lake provides ACID transactions, scalable metadata handling, and unifies streaming and batch data processing. Delta Lake runs on top of your existing data lake and is fully compatible with Apache Spark APIs.

Specifically, Delta Lake offers:

- ACID (Atomicity, Concurrency, Isolation, Durability) transactions on Data Lake: Serializable isolation levels ensure that readers never see inconsistent data.
- Scalable metadata handling: Leverages Spark's distributed processing power to handle all the metadata for petabyte-scale tables with billions of files at ease.
- Streaming and batch unification: A table in Delta Lake can be a batch as well as a streaming source and sink. Streaming data ingest, batch historic backfill, interactive queries all just work out of the box.
- Schema enforcement: Automatically handles schema variations to prevent insertion of bad records during ingestion.
- Time travel: Data versioning enables rollbacks, full historical audit trails, and reproducible machine learning experiments.
- Upserts and deletes: Supports merge, update and delete operations to enable complex use cases like change-data-capture, slowly-changing-dimension (SCD) operations, streaming upserts, and so on.

<https://docs.delta.io/latest/delta-intro.html>

## Deltalake (cont...)

- Delta Lake is open source software that extends Parquet data files with a file-based transaction log for ACID transactions and scalable metadata handling.
- Delta Lake is the default format for all operations on Azure Databricks. Unless otherwise specified, all tables on Azure Databricks are Delta tables.
- Databricks originally developed the Delta Lake protocol and continues to actively contribute to the open source project.

<https://learn.microsoft.com/en-us/azure/databricks/delta/tutorial>

# Databricks Delta Lake Overview

- Key Features:
  - ✓ ACID transactions (multiple concurrent reads and writes)
  - ✓ Open format (Parquet)
  - ✓ Scalable metadata handling
  - ✓ Streaming and batch unification
  - ✓ Schema enforcement
  - ✓ Time travel
  - ✓ Upserts and deletes
  - ✓ Audit history
  - ✓ 100% compatible with Apache Spark API

## Key Features

### ACID Transactions:

Data lakes typically have multiple data pipelines reading and writing data concurrently, and data engineers have to go through a tedious process to ensure data integrity, due to the lack of transactions. Delta Lake brings ACID transactions to your data lakes. It provides serializability, the strongest level of isolation level. Learn more at [Diving into Delta Lake: Unpacking the Transaction Log](#).

### Scalable Metadata Handling:

In big data, even the metadata itself can be "big data". Delta Lake treats metadata just like data, leveraging Spark's distributed processing power to handle all its metadata. As a result, Delta Lake can handle petabyte-scale tables with billions of partitions and files at ease.

### Time Travel (data versioning):

Delta Lake provides snapshots of data enabling developers to access and revert to earlier versions of data for audits, rollbacks or to reproduce experiments. Learn more in [Introducing Delta Lake Time Travel for Large Scale Data Lakes](#).

### Open Format:

All data in Delta Lake is stored in Apache Parquet format enabling Delta Lake to leverage the efficient compression and encoding schemes that are native to Parquet.

### Unified Batch and Streaming Source and Sink:

A table in Delta Lake is both a batch table, as well as a streaming source and sink.

Streaming data ingest, batch historic backfill, and interactive queries all just work out of the box.

#### Schema Enforcement:

Delta Lake provides the ability to specify your schema and enforce it. This helps ensure that the data types are correct and required columns are present, preventing bad data from causing data corruption. For more information, refer to [Diving Into Delta Lake: Schema Enforcement & Evolution](#).

#### Schema Evolution:

Big data is continuously changing. Delta Lake enables you to make changes to a table schema that can be applied automatically, without the need for cumbersome DDL. For more information, refer to [Diving Into Delta Lake: Schema Enforcement & Evolution](#).

#### Audit History:

Delta Lake transaction log records details about every change made to data providing a full audit trail of the changes.

#### Updates and Deletes:

Delta Lake supports Scala / Java APIs to merge, update and delete datasets. This allows you to easily comply with GDPR and CCPA and also simplifies use cases like change data capture. For more information, refer to [Announcing the Delta Lake 0.3.0 Release](#) and [Simple, Reliable Upserts and Deletes on Delta Lake Tables using Python APIs](#) which includes code snippets for merge, update, and delete DML commands.

#### 100% Compatible with Apache Spark API:

Developers can use Delta Lake with their existing data pipelines with minimal change as it is fully compatible with Spark, the commonly used big data processing engine.

<https://delta.io/>

## Databricks Delta Lake Overview

- Delta Lake uses versioned Parquet files to store your data in your cloud storage.
- Delta Lake also stores a transaction log to keep track of all the commits made to the table or blob store directory to provide ACID transactions



Refer to these links for review material related to Databricks Delta. Basically, Databricks Delta optimizes storage of data in the Databricks File System through the use of Delta tables which use the Parquet format as their underlying file format.

<https://docs.azuredatabricks.net/delta/delta-intro.html#delta-intro>

<https://docs.azuredatabricks.net/delta/optimizations.html>

# Databricks Delta Lake Overview

- Entries in the log, called *delta files*, are stored as atomic collections of [actions](#) in the `_delta_log` directory, at the root of a table
- Entries in the log are encoded using JSON and are named as zero-padded contiguous integers.

The screenshot shows the Databricks Data Explorer interface with three main sections:

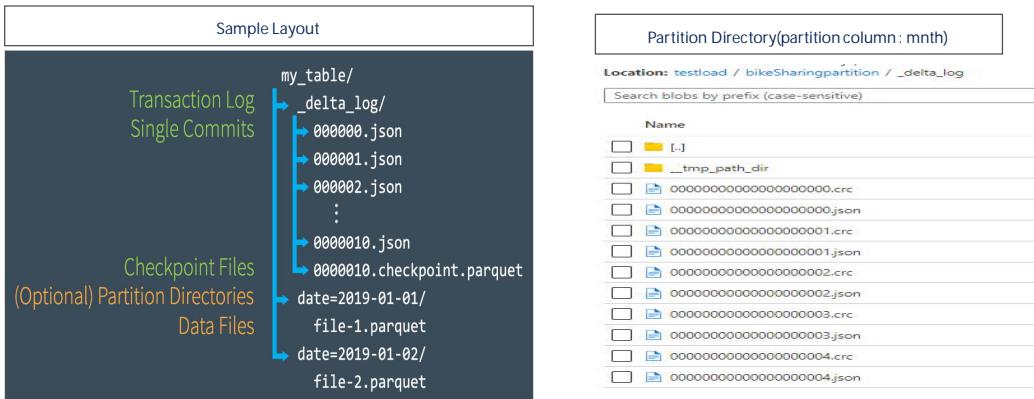
- \_delta\_log directory**: Located at `testload / bikeSharingpartition / _delta_log`. It contains entries for `[-]`, `_tmp_path_dir`, and two JSON files: `00000000000000000000000000000000.crc` and `00000000000000000000000000000000.json`.
- Partition Directory(partition column : mnth)**: Located at `testload / bikeSharingpartition / mnth=1`. It contains entries for `[-]` and a single Parquet file: `part-00000-fbec9d52-28d2-4b89-a11c-5d0db15e9ffb.c000.snappy.parquet`.
- Partition Folder layout Example(mnth is a partition column)**: A callout box showing the contents of the `mnth=1` partition. It lists integer values from 1 to 12, each preceded by a yellow folder icon.

<https://databricks.com/blog/2019/08/21/diving-into-delta-lake-unpacking-the-transaction-log.html>

Walk through pictorial representation of `_delta_log` file, partition folder contents and partition root folder layout.

# Databricks Delta Lake overview

- Delta Lake automatically generates checkpoint files every 10 commits
- These checkpoint files save the entire state of the table at a point in time – in native Parquet format that is quick and easy for Spark to read



<https://databricks.com/blog/2019/08/21/diving-into-delta-lake-unpacking-the-transaction-log.html>

## Deltalake (cont...)

```
1 %sql
2 CREATE OR REPLACE TABLE scd2demo(
3   pk1 INT,
4   pk2 STRING,
5   dim1 INT,
6   dim2 int,
7   dim3 int,
8   dim4 int,
9   active_status STRING,
10  start_date TIMESTAMP,
11  end_date TIMESTAMP)
12 LOCATION 'FileStore/tables/scd2demo'
```

- ▶ (4) Spark Jobs
- ▶ \_sqldf: pyspark.sql.dataframe.DataFrame

OK

/FileStore/tables/scd2demo/\_delta\_log

Q Prefix search

\_delta\_log

Q Prefix search

\_tbl\_path\_dir

00000000000000000000000000000000.crc

00000000000000000000000000000000.json

## Deltalake (cont...)

```
%sql  
insert into scd2demo values(111,'Unit1',200,300,400,500,'Y',current_timestamp(),'9999-12-31');  
insert into scd2demo values(222,'Unit2',200,300,400,500,'Y',current_timestamp(),'9999-12-31');  
insert into scd2demo values(333,'Unit3',200,300,400,500,'Y',current_timestamp(),'9999-12-31');
```

The screenshot shows the Databricks File Store interface. On the left, there is a sidebar with a search bar labeled 'Prefix search' and a dropdown menu showing 'scd2demo'. On the right, there is another search bar labeled 'Prefix search' and a list of files under '\_delta\_log'. The list contains three files: 'part-00000-03e48024-dc93-46db-...', 'part-00000-63352cea-c487-4bdb-...', and 'part-00000-c85a16ae-1d24-42f4-b...'. A red box highlights the text 'These are data files for our table' above the file list, and a red arrow points from this box to the files.

<https://medium.com/@amarkrgupta96/unlocking-performance-optimize-vacuum-and-z-ordering-in-databricks-delta-tables>

# Table Batch Reads and Writes

Delta Lake supports most of the options provided by a DataFrame read and write APIs for performing batch reads and writes on tables

Spark SQL	Dataframe API in Pyspark
<b>Create a table</b> CREATE TABLE bi keShari ng USING DELTA LOCATION '/tmp/bi keShari ng/'	<b>Create a table</b> df = spark.read.format("delta").load("/tmp/delta-table") df.write.format("delta").save("/tmp/delta-table")
<b>Partition Data</b> CREATE TABLE bi keShari ngpartition USING DELTA PARTITIONED BY (mnth) LOCATION '/tmp/bi keShari ng/'	<b>Partition Data</b> df = spark.read.format("delta").load("/tmp/delta-table") df.write.format("delta").partitionBy("date").save("/tmp/delta-table")
<b>Read Data</b> SELECT * FROM bi keShari ng	<b>Read Data</b> df = spark.read.format("delta").load("/tmp/delta-table") df.show()

<https://docs.delta.io/latest/delta-batch.html>

# Table Batch Reads and Writes

Read older versions of data using time travel

- As you write into a Delta table or directory, every operation is automatically versioned. You can access the different versions of the data two different ways via version number or TimeStamp

Using Version

```
SELECT count(1) from bikeSharingDay VERSION AS OF 1  
SELECT count(1) from bikeSharingDay VERSION AS OF 0
```

Using timestamp

```
SELECT * FROM bikeSharingDay TIMESTAMP AS OF '2019-01-29 00:37:58'  
SELECT * FROM bikeSharingpartition TIMESTAMP AS OF '2019-01-29 00:37:58'
```

Introduction to time travel

<https://databricks.com/blog/2019/02/04/introducing-delta-time-travel-for-large-scale-data-lakes.html>

- timestamp\_expression can be any one of:
  - '2018-10-18T22:15:12.013Z', that is, a string that can be cast to a timestamp
  - cast('2018-10-18 13:36:32 CEST' as timestamp)
  - '2018-10-18', that is, a date string
  - current\_timestamp() - interval 12 hours
  - date\_sub(current\_date(), 1)
  - Any other expression that is or can be cast to a timestamp
- version is a long value that can be obtained from the output of DESCRIBE HISTORY table\_spec.

Neither timestamp\_expression nor version can be subqueries.

## Data retention

- To time travel, you must retain both the log and the data files for that version
    - The data files are never deleted automatically. You need to run VACUUM for this to happen. Vacuum does not delete log files. Log files are automatically cleaned up after checkpoints are written
    - By default you can time travel up to 30 days unless you have
      - Run Vacuum
      - Changed the data or log retention periods by using the following table.properties
        - delta.logRetentionDuration = "interval <interval>" (default 30 days)
        - delta.deletedFileRetentionDuration = "interval <interval>" (default 7 days)
- In order to have 30 days full history, you'd have to have delta.deletedFileRetentionDuration set to "interval 30 days"

# Table Batch Reads and Writes

## Read older versions of data using time travel

You can look at the history of table changes using the DESCRIBE HISTORY command or through the UI – the Data Explorer (classic schema browser)

List change history on the table

DESCRIBE HISTORY bi\_keShari\_ngDay

The image contains two side-by-side screenshots of the Databricks Data Explorer interface. The left screenshot shows a list of tables in the 'default' database, with 'patient' being the only listed table. The right screenshot shows the same view, but with the 'History' tab selected for the 'bi\_keShari\_ngDay' table. This tab displays a table of historical data changes, with one row visible:

Version	Timestamp	User Id	User Name	Operation
0	2023-05-09T17:51:46	0000111111111111	chris@mringen209591.onmicrosoft.com	CREATE TABLE AS SELECT

## Introduction to time travel

<https://databricks.com/blog/2019/02/04/introducing-delta-time-travel-for-large-scale-data-lakes.html>

Be aware, if you wanted to demo the UI, you need to have a warehouse started. (you switch view to SQL, then SQL Warehouses and then either start the starter warehouse or create a new one)

# Table Batch Reads and Writes

## Update a Table

Spark SQL	Dataframe API in Pyspark
<b>Append</b> <pre>INSERT INTO bikesharing SELECT * FROM bikesharing_temp</pre>	<b>Append</b> <pre>df.write.format("delta").mode("append").save("/tmp/delta-table")</pre>
<b>Overwrite</b> <pre>INSERT OVERWRITE TABLE bikesharing SELECT * FROM bikesharing_temp</pre>	<b>Overwrite</b> <pre>df.write.format("delta").mode("overwrite").save("/tmp/delta-table")</pre>
You can selectively overwrite only the data that matches predicates over partition columns	You can selectively overwrite only the data that matches predicates over partition columns <pre>df.write   .format("delta")   .mode("overwrite")   .option("replaceWhere", "date &gt;= '2020-01-01' AND           date &lt;= '2020-01-31' ")   .save("/tmp/delta-table")</pre>
<b>UPDATE</b> bikesharing <pre>SET cnt = cnt+1 where dteday between "2011-01-01" and "2011-01-01"</pre>	

<https://docs.delta.io/latest/delta-batch.html>

# Table Batch Reads and Writes

## Update a Table – Fix incorrect updates

Spark SQL	Dataframe API in Pyspark
<p>Fix the cnt column and increment it by one where date falls between "2011-01-01" and "2011-01-01"</p> <pre>UPDATE bikesharing SET cnt = cnt+1 where dteday between "2011-01-01" and "2011-01-01"</pre>	<p>You can update data that matches a predicate in a Delta table. For example, to fix a spelling mistake in the eventType, you can run the following:</p> <pre>from delta.tables import * from pyspark.sql.functions import * deltaTable = DeltaTable.forName(spark, "/data/events/") deltaTable.update("eventType = 'clk'", { "eventType": "click" }) ) # predicate using SQL formatted string /* deltaTable.update(col("eventType") == "clk", { "eventType": lit("click") }) # predicate using Spark SQL functions */</pre>

<https://docs.delta.io/latest/delta-batch.html>

# Table Batch Reads and Writes

## Update a Table - Upsert

Spark SQL	Dataframe API in Pyspark
<p>Upsert (Merge)</p> <pre>MERGE INTO my_table target USING my_table TIMESTAMP AS OF date_sub(current_date(), 1) source ON source.userId = target.userId WHEN MATCHED THEN UPDATE SET *</pre>	<p>Upsert (Merge)</p> <p>You can upsert data from a DataFrame into a Delta table using the merge operation</p> <pre>from delta.tables import * deltaTable = DeltaTable.forPath(spark, "/data/events/") deltaTable.alias("events").merge(updatesDF.alias("updates"), "events.eventId = updates.eventId") \ .whenMatchedUpdate(set = { "data" : "updates.data" }) \ .whenNotMatchedInsert(values = { "date": "updates.date", "eventId": "updates.eventId", "data": "updates.data" }) \ .execute()</pre>

<https://docs.delta.io/latest/delta-update.html#-delta-update>

<https://docs.databricks.com/delta/delta-batch.html#language-sql>

# Table Batch Reads and Writes

## Update a Table

- Delete

You can remove data that matches a predicate from a Delta table

Spark SQL	Dataframe API in Pyspark
<pre>DELETE FROM events WHERE date &lt; '2017-01-01' DELETE FROM delta.`/data/events` WHERE date &lt; '2017-01-01'</pre>	<pre>from delta.tables import * from pyspark.sql.functions import * deltaTable = DeltaTable.forPath(spark, "/data/events/") deltaTable.delete("date &lt; '2020-01-01'") # # predicate using SQL formatted string /* deltaTable.delete(col ("date") &lt; "2020-01-01") # predicate using Spark SQL functions */</pre>

<https://docs.delta.io/latest/delta-update.html#-delta-update>

## Other Capabilities

- Dynamic Partition Overwrites
- Limit Rows written to one file
- Idempotent Writes (appid:version)
- Schema Validation
- Update Table Schema
  - Explicit via Alter Table Add Columns
  - Change column comment or ordering
  - Replace Columns
  - Rename Columns
  - Drop Columns
  - Change Column Type or Name

- Dynamic partition overwrites
- When in dynamic partition overwrite mode, we overwrite all existing data in each logical partition for which the write will commit new data. Any existing logical partitions for which the write does not contain data will remain unchanged. This mode is only applicable when data is being written in overwrite mode: either INSERT OVERWRITE in SQL, or a DataFrame write with `df.write.mode("overwrite")`.
- Configure dynamic partition overwrite mode by setting the Spark session configuration `spark.sql.sources.partitionOverwriteMode` to `dynamic`. You can also enable this by setting the `DataFrameWriter` option `partitionOverwriteMode` to `dynamic`. If present, the query-specific option overrides the mode defined in the session configuration. The default for `partitionOverwriteMode` is `static`.

### Limit rows written to one file

You can use the SQL session configuration `spark.sql.files.maxRecordsPerFile` to specify the maximum number of records to write to a single file for a Delta Lake table. Specifying a value of zero or a negative value represents no limit.

You can also use the `DataFrameWriter` option `maxRecordsPerFile` when using the `DataFrame` APIs to write to a Delta Lake table. When `maxRecordsPerFile` is specified, the value of the SQL session configuration `spark.sql.files.maxRecordsPerFile` is ignored.

## [Idempotent writes](#)

Sometimes a job that writes data to a Delta table is restarted due to various reasons (for example, job encounters a failure). The failed job may or may not have written the data to Delta table before terminating. In the case where the data is written to the Delta table, the restarted job writes the same data to the Delta table which results in duplicate data.

To address this, Delta tables support the following DataFrameWriter options to make the writes idempotent:

- `txnAppId`: A unique string that you can pass on each DataFrame write. For example, this can be the name of the job.
- `txnVersion`: A monotonically increasing number that acts as transaction version. This number needs to be unique for data that is being written to the Delta table(s). For example, this can be the epoch seconds of the instant when the query is attempted for the first time. Any subsequent restarts of the same job needs to have the same value for `txnVersion`.

The above combination of options needs to be unique for each new data that is being ingested into the Delta table and the `txnVersion` needs to be higher than the last data that was ingested into the Delta table. For example:

- Last successfully written data contains option values as `dailyETL:23423 (txnAppId:txnVersion)`.
- Next write of data should have `txnAppId = dailyETL` and `txnVersion` as at least 23424 (one more than the last written data `txnVersion`).
- Any attempt to write data with `txnAppId = dailyETL` and `txnVersion` as 23422 or less is ignored because the `txnVersion` is less than the last recorded `txnVersion` in the table.
- Attempt to write data with `txnAppId:txnVersion` as `anotherETL:23424` is successful writing data to the table as it contains a different `txnAppId` compared to the same option value in last ingested data.

You can also configure idempotent writes by setting the Spark session configuration `spark.databricks.delta.write.txnAppId` and `spark.databricks.delta.write.txnVersion`. In addition, you can set `spark.databricks.delta.write.txnVersion.autoReset.enabled` to true to automatically reset `spark.databricks.delta.write.txnVersion` after every write. When both the writer options and session configuration are set, we will use the writer option values.

### Warning

This solution assumes that the data being written to Delta table(s) in multiple retries of the job is same. If a write attempt in a Delta table succeeds but due to some downstream failure there is a second write attempt with same txm options but different data, then that second write attempt will be ignored. This can cause unexpected results.

## [Schema validation](#)

Delta Lake automatically validates that the schema of the DataFrame being written is compatible with the schema of the table. Delta Lake uses the following rules to determine whether a write from a DataFrame to a table is compatible:

- All DataFrame columns must exist in the target table. If there are columns in the

DataFrame not present in the table, an exception is raised. Columns present in the table but not in the DataFrame are set to null.

- DataFrame column data types must match the column data types in the target table. If they don't match, an exception is raised.
- DataFrame column names cannot differ only by case. This means that you cannot have columns such as "Foo" and "foo" defined in the same table. While you can use Spark in case sensitive or insensitive (default) mode, Parquet is case sensitive when storing and returning column information. Delta Lake is case-preserving but insensitive when storing the schema and has this restriction to avoid potential mistakes, data corruption, or loss issues.

Delta Lake support DDL to add new columns explicitly and the ability to update schema automatically.

If you specify other options, such as partitionBy, in combination with append mode, Delta Lake validates that they match and throws an error for any mismatch. When partitionBy is not present, appends automatically follow the partitioning of the existing data.

- Update Table Schema
- <https://docs.delta.io/2.3.0/delta-batch.html#id25>

## Change Data Feed

- Tracks row level changes between versions of a Delta table
- 'Change events' are recorded by runtime for the data written to the table
  - These events consist of data plus metadata whether data was inserted, updated or deleted
- Silver To Gold, Transmit Changes, Audit Trail
- Enable by
  - New table
    - Create TABLE.... TBLPROPERTIES (delta.enableChangeDataFeed = true)
  - Existing Table
    - ALTER TABLE <table> SET TBLPROPERTIES (delta.enableChangeDataFeed = true)
  - All new Tables
    - Set spark.databricks.delta.properties.defaults.enableChangeDataFeed = true

See

<https://docs.delta.io/2.3.0/delta-change-data-feed.html#change-data-feed>

## Optimize performances with file management

- To improve query speed, Delta Lake offers the ability to optimize the layout of data stored in storage
- Delta Lake supports two layout algorithms: **bin-packing** and **Z-Ordering**

### Compaction (bin-packing)

- Bin-packing coalesces small files into larger ones, avoiding scanning many files when executing lookup queries
- You trigger compaction by running the OPTIMIZE command

`OPTIMIZE delta.`/data/events``

- If you have a large amount of data and only want to optimize a subset of it, you can specify an optional partition predicate using WHERE

`OPTIMIZE delta.`/data/events` WHERE date >= '2017-01-01'`

<https://databricks.com/blog/2018/07/31/processing-petabytes-of-data-in-seconds-with-databricks-delta.html>

# OPTIMIZE

```
%sql  
OPTIMIZE scd2demo;  
  
+---+  
| (8) Spark Jobs  
| _sqlId: pyspark.sql.dataframe.DataFrame = [path: string, metrics: struct]  
+---+  
| Table +  
| path | metrics  
+---+  
| dbfs:/FileStore/tables/scd2demo | {"numFilesAdded": 1, "numFilesRemoved": 5, "filesAdded": [{"min": 2333, "max": 2333, "avg": 2333, "totalFiles": 1, "totalSize": 2333}, {"filesRemoved": [{"min": 2033, "max": 2073, "avg": 2033.4, "totalFiles": 5, "totalSize": 10167}], "partitionsOptimized": 0, "zOrderStats": null, "numBatches": 1, "totalConsideredFiles": 5, "totalFilesSkipped": 0, "preserveInsertionOrder": true, "numFilesSkippedToReduceWriteAmplification": 0, "numBytesSkippedToReduceWriteAmplification": 0, "startTimeMs": 1685894350328, "endTimeMs": 1685894353574, "totalClusterParallelism": 16, "totalScheduledTasks": 1, "autoCompactParallelismStats": null}]  
+---+
```

## Optimize performances with file management- Continued

### Data skipping

- As new data is inserted into a Delta table, file-level **min/max** statistics are collected for all columns of supported types
- When there's a lookup query against the table, Delta table first consults these statistics in order to determine which files can safely be skipped
- You do not need to configure data skipping - the feature is activated whenever applicable. However, its effectiveness depends on the layout of your data. For best results, apply [Z-Ordering](#).

<https://databricks.com/blog/2018/07/31/processing-petabytes-of-data-in-seconds-with-databricks-delta.html>

# Optimize performances with file management

## Z-Ordering

- Z-Ordering is a [technique](#) to colocate related information in the same set of files
- This co-locality is automatically used by data-skipping algorithms to dramatically reduce the amount of data that needs to be read
- To Z-Order data, you specify the columns to order on in the ZORDER BY clause:

```
%sql  
OPTIMIZE events WHERE date >= current_timestamp() - INTERVAL 1 day  
ZORDER BY (eventType)
```

<https://databricks.com/blog/2018/07/31/processing-petabytes-of-data-in-seconds-with-databricks-delta.html>

# Optimize performances with file management

## Z-Ordering

- If you expect a column to be commonly used in query predicates and if that column has high cardinality, then use ZORDER BY
- You can specify multiple columns for ZORDER BY as a comma-separated list. However, the effectiveness of the locality drops with each additional column

<https://databricks.com/blog/2018/07/31/processing-petabytes-of-data-in-seconds-with-databricks-delta.html>

**Employee Delta Table**

File 1				File 3				File 5			
<b>Emp_id</b>	<b>Emp_name</b>	<b>Salary</b>	<b>Active</b>	<b>Emp_id</b>	<b>Emp_name</b>	<b>Salary</b>	<b>Active</b>	<b>Emp_id</b>	<b>Emp_name</b>	<b>Salary</b>	<b>Active</b>
111 Michael	5000 Y	222 Nancy	5000 Y	555 Kevin	3000 N						
555 Kevin	7000 Y	444 Tomas	6000 Y	888 Shane	4500 Y						

**Optimized Employee Table**

File 1				File 2				File 3			
<b>Emp_id</b>	<b>Emp_name</b>	<b>Salary</b>	<b>Active</b>	<b>Emp_id</b>	<b>Emp_name</b>	<b>Salary</b>	<b>Active</b>	<b>Emp_id</b>	<b>Emp_name</b>	<b>Salary</b>	<b>Active</b>
333 David	4000 Y	222 Nancy	5000 Y	555 Kevin	3000 N	444 Tomas	6000 Y	888 Shane	4500 Y		
999 Peter	3000 Y	333 David	2000 N	111 Michael	5000 Y	444 Tomas	6000 Y	444 Tomas	2500 N	999 Peter	2000 N

select * from employee where emp_id =444			select * from employee where emp_id <333			File	Min	Max
1,2,3	1,2	1	File1	111	999			
			File2	222	888			
			File3	444	999			

2
%sql

3
OPTIMIZE scd2Demo

ZORDER PK1

**Z-Ordered Employee Table**

File 1				File 2				File 3			
<b>Emp_id</b>	<b>Emp_name</b>	<b>Salary</b>	<b>Active</b>	<b>Emp_id</b>	<b>Emp_name</b>	<b>Salary</b>	<b>Active</b>	<b>Emp_id</b>	<b>Emp_name</b>	<b>Salary</b>	<b>Active</b>
111 Michael	5000 Y	444 Tomas	6000 Y	222 Nancy	5000 Y	555 Kevin	7000 Y	333 David	4000 Y	555 Kevin	3000 N
222 Nancy	5000 Y	444 Tomas	2500 N	333 David	4000 Y	555 Kevin	7000 Y	111 Michael	5000 Y	555 Kevin	3000 N
333 David	4000 Y	555 Kevin	7000 Y	333 David	2000 N	555 Kevin	3000 N				
333 David	2000 N	555 Kevin	3000 N								

select * from employee where emp_id =444			select * from employee where emp_id <333			File	Min	Max
2	1	1	File1	111	333			
			File2	444	555			
			File3	888	999			

<https://medium.com/@amarkrgupta96/unlocking-performance-optimize-vacuum-and-z-ordering-in-databricks-delta-tables>

# Table Utility commands

## Vaccum

- Remove files no longer referenced by a Delta table and are older than the retention threshold
- The default retention threshold for the files is 7 days.
- Vacuum is not triggered automatically.

```
VACUUM [db_name.]table_name|path [RETAIN num HOURS] [DRY RUN]
```

%sql

```
VACUUM delta.`/data/events/` -- vacuum files not required by versions older than the default retention period
```

```
VACUUM delta.`/data/events/` RETAIN 100 HOURS -- vacuum files not required by versions more than 100 hours old
```

```
VACUUM delta.`/data/events/` DRY RUN-- Return a list of files to be deleted
```

<https://docs.databricks.com/spark/latest/spark-sql/language-manual/vacuum.html>

By default, retain hours for vacuum should be greater than 168 hour, which can be altered to a lower number with following setting :

```
spark.databricks.delta.retentionDurationCheck.enabled to false
```

## Table Utility commands - Contd

```
DESCRIBE HISTORY '/data/events/' -- get the full history of the table
DESCRIBE HISTORY '/data/events/' LIMIT 1 -- get the last operation only

DESCRIBE DETAIL <DATABASENAME.TABLENAME> -- Show table details

SHOW DATABASES -- List Databases
SHOW TABLES -- List tables
SHOW CREATE TABLE <DATABASENAME.TABLENAME> -- Show table creation script

RESTORE TABLE <DATABASENAME.TABLENAME> TO VERSION AS OF <version>
RESTORE TABLE '/data/target/' TO TIMESTAMP AS OF <timestamp>
```

<https://docs.delta.io/latest/delta-utility.html>

## Demo: Databricks Delta Table Demo

Exercise 1: Create and Modify Delta table  
Exercise 2: Create Managed Parquet tables  
Exercise 3: Optimize and Analyze Delta Table  
Exercise 4 : Data Skipping Technique  
Exercise 5 : Time Travel



Refer to M03\_L02\_Demo02.ipynb file for Databricks notebook code. Remember to import the notebook to your workspace and read the directions before executing the code

## Knowledge Check

1. How can you optimize Delta lake performances
2. What is the impact of VACUUM table utility command

1. How can you optimize Delta lake performances?

- By using Bin-Packing and Zordering

2. What is the impact of VACUUM table utility command

Recursively vacuum directories associated with the Delta table and remove files that are no longer in the latest state of the transaction log for the table and are older than a retention threshold

## Module Summary

- In this module we covered the following:
  - Delta Lake is a storage layer that brings scalable, ACID transactions to [Apache Spark](#) and other big-data engines
  - Data Skipping Technique of Delta Table
  - Optimize and Analyze Delta Table
  - Time Traveler capability