





Feature Engineering with Databricks

CS5304

Data Science & AI on Databricks






Features of Databricks for Machine Learning


Non exhaustive list of features that will be used throughout this module

- Collaborative notebooks
- ML Runtime
- Governance of Data & Models (*via Unity Catalog*)
- Feature Store
- Managed MLflow
- Model Serving
- AutoML



Quick Exploratory Data Analysis

Native tools for visualizing and understanding data in ML workflow



Create **interactive charts** to visualize data in the Notebook with only two clicks



Summarize a data set's essential properties and statistics in a **data profile** with the push of a button

©2017 Databricks Inc. - All rights reserved



Data Preparation for ML projects

Goal: Optimize input quality for accurate model predictions

Data preparation includes the following tasks:


- **Cleaning and formatting data:** This includes tasks such as **handling missing values** or outliers, ensuring data is in the correct format, and removing unneeded columns.
- **Feature Engineering:** This includes tasks like numerical **transformations**, **aggregating data**, encoding text or image data, and **creating new features**.



Feature Extraction

Transforming raw data into a set of features that better represent the underlying patterns


- Transforming Raw Data for Enhanced Modeling
- Dimensionality Reduction for Improved Performance
- Simplifying Feature Engineering



Feature Selection: Dropping Columns/Rows


Enhancing Model Efficiency

Streamlining the dataset by selectively removing columns (features) or rows (instances) to improve model efficiency and effectiveness.



Why Do We Need **Splitting Data**

Optimizing Model Evaluation



©2025 Databricks Inc. — All rights reserved.




k

Splitting Data into Multiple Sets

Optimizing Model Training, Validation, and Testing


Splitting the dataset is a fundamental step in data preparation to facilitate effective model training, validation, and testing.



Data Imputation

Data imputation is **the process of filling in missing values** in a dataset with estimated or predicted values.

The goal of data imputation is to enhance the quality and completeness of the dataset, ultimately improving the performance and reliability of the machine learning model.



Problems with Missing Data

Impacting the performance and reliability of ML models

- Reduced Model Performance
- Biased Inferences
- Imbalanced Representations
- Increased Complexity in Model


Handling

Column	Data type	Unique values	Missing values
Pclass	int64	3	0
Survived	int64	2	0
Name	int64	3	0
Sex	object	2	0
Age	float64	80	177 - 10.3%
SibSp	int64	7	0
Parch	int64	7	0
Ticket	object	381	0
Fare	float64	246	0
Cabin	object	140	687 - 77.1%
Embarked	object	4	2 - 0.2%



How to Handle Missing Data

Data imputation methods



Replacing Missing Values

Data imputation methods

Mean - Mode Imputation

Before	After
10	10.0
15	15.0
-	18.3
20	20.0
25	25.0

K-Nearest Neighbors
(KNN with k=2)

Before	After
8	8.0
-	10.0
12	12.0
15	15.0
-	13.0

Multiple Imputation
(Regression)

F1	F2	Before	After
X	X	10	10.0
X	X	15	15.0
● X	$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \epsilon$		Y
X	X	20	20.0
X	X	25	25.0

©2025 Databricks Inc. – All rights reserved.



Marking Imputed Data (*Best Practice*)

Keep track of imputed data

Important to mark imputed data, for:

- Model Evaluation
- Data Quality Assessment
- Enabling Transparency of Dataset
- Error Identification

ID	Name	Age	Age_imputed
1	Alice	25.0	0
2	Bob	30.0	0
3	Charlie	26.0	1
4	David	28.0	0
5	Eva	22.0	0

Data Encoding



Why Encoding?

Data encoding is an important pre-processing step in preparing categorical data for machine learning algorithms, as vast majority of algorithms accept numerical input exclusively.

Issues with classic ML

- Handling high-cardinality features
- Introducing unintended relationships
- Overfitting
- Increased computational cost
- Possible lack of interpretability

Encoding Process



Working with Categorical Features

High Cardinality

Issue: Large set of categories

Many unique values in a categorical feature can lead to a large number of dummy variables, increasing dimensionality and potentially causing issues. This is known as the **high-cardinality problem**.

Possible Solutions:

Group Rare Categories:

ID	Category
0	A
1	B
2	A
3	C
4	D
5	D
6	D
7	D

ID	Category	Cat_Group
0	A	A
1	B	Rare
2	A	A
3	C	Rare
4	D	D
5	D	D
6	D	D
7	D	D

Top-N Categories:

ID	Category	Cat_Group
0	A	A
1	B	Other
2	A	A
3	C	Other
4	D	D
5	D	D
6	D	D
7	D	D

©2025 Databricks Inc. All rights reserved.

Working with Categorical Features

Missing Values

Issue: Categorical gaps

Categorical features often have **missing values**, which need to be addressed before model training.

Possible Solutions:

Imputation:

ID	Feature A	Feature B
0	1.0	10.0
1	2.0	NaN
2	NaN	30.0
3	4.0	40.0
4	5.0	50.0

Consider Missing as a Separate Category:

ID	Feature A	Feature B
0	1.0	10.0
1	2.0	NaN
2	NaN	30.0
3	4.0	40.0
4	5.0	50.0

Working with Categorical Features

Encoding Categories

Issue: String types

Models require numerical input, and categorical variables need to be **encoded**.

Possible Solutions:

One-Hot Encoding

ID	Category	ID	Cat_A	Cat_B	Cat_C
0	A	0	1	1	1
1	B	1	0	0	0
2	A	2	1	1	1
3	C	3	0	0	0
4	A	4	1	1	1

Label Encoding:

ID	Category	ID	Category	Label
0	A	0	A	0
1	B	1	B	1
2	A	2	A	0
3	C	3	C	2
4	A	4	A	0

Ordinal Encoding:

ID	Category	ID	Category	Ordinal
0	A	0	A	0.0
1	B	1	B	1.0
2	A	2	A	0.0
3	C	3	C	2.0
4	A	4	A	0.0

Label Encoding for Ordinal Features

Convert categorical data into numerical labels (aka "String Indexing")

Procedure:

1. **Assign numeric labels:** Map each category to a numeric value based on its natural order.
2. **Transform the Feature:** Replace each categorical value in the feature column with its corresponding numeric label.


Example:

String Value	Numeric Value
Freshman	1
Sophomore	2
Junior	3
Senior	4

The diagram illustrates the process of label encoding for ordinal features. It shows two tables. The top table, titled 'High School Grade Level', maps categorical values to numeric values: Freshman (1), Senior (2), Junior (3), and Sophomore (4). An arrow points from this table to a second table below, which is a copy of the first but with the 'High School Grade Level' column replaced by its corresponding numeric values (1, 2, 3, 4). A large downward-pointing arrow is positioned between the two tables, indicating the transformation.

ID	High School Grade Level	Age
1	Freshman	14
2	Senior	17
3	Junior	16
4	Freshman	15
5	Sophomore	16


ID	High School Grade Level	Age
1	1	14
2	4	17
3	3	16
4	1	15
5	2	16



Feature Store

Data Preparation for Machine Learning

©2025 Databricks Inc. — All rights reserved



What is a Feature Store?

A **feature store** manages features, or input data to a machine learning model.

In a model that predicts **customer churn**, for example, features could be:

- Aggregations of raw data over time windows, like **trailing 7-day purchases**
- Joined **combinations of data sets**, like customer demographic information joined to transaction features
- Complex functions of customer information, like **estimated customer lifetime value**

The process of creating these values from data is **feature engineering**.

Why Would You Need a Feature Store?

Basic Motivations

Discovery

Multiple Data Scientists are trying to solve similar modeling tasks and come up with different definitions of the same features.
How can I find the features?

Lineage

Model governance requires documentation of the features used to train a model, as well as the **upstream lineage** of a feature to reliably use it. **How is it computed, and who owns it?**

Skew

When multiple teams manage feature computation and ML models in production, minor yet significant **skew in upstream data** at the input of a feature pipeline can be very hard to detect and fix.

Online Serving


During exploration and model experimentation phases features are implemented in frameworks that do not scale to production.




What is a Feature Store?

An example of a recommendation system.

A **feature store** manages features, or input data to a machine learning model.



Databricks Feature Store



Complete Integration-FS with Unity Catalog

Any table can be a feature table

- Feature Tables become regular UC Tables with additional metadata.
- Shared properties are unified.
 - Feature table description == table comment.
 - Feature table schema == table schema
- Three-level namespace convention

The screenshot shows the Databricks Data Explorer interface. On the left, there is a sidebar with navigation options like Home, Databricks, Workspaces, and Databricks Catalog. The main area displays a table named 'LOG_JAVAFX'. The table has the following columns:

Name	Type	Desc	Comment
id	string		
name	string		
type	string		
desc	string		
comment	string		

©2025 Databricks Inc. — All rights reserved.





Feature Engineering Demo

Load Dataset

```
# Load dataset with spark
shared_volume_name = 'telco' # From Marketplace
csv_name = 'telco-customer-churn' # CSV file name
dataset_path = f'{DA.paths.datasets.telco}/{shared_volume_name}/{csv_name}.csv' # Full path
telco_df = spark.read.csv(dataset_path, header="true", inferSchema="true", multiLine="true", escape="''")

# # Drop the target column
telco_df = telco_df.drop("Churn")

# # View dataset
display(telco_df)
```

Create Feature Table

```
# # create a feature table from the dataset
table_name = f"{DA.catalog_name}.{DA.schema_name}.telco_customer_features"

fe.create_table(
    name=table_name,
    primary_keys=["customerID"],
    df=telco_df,
    #partition_columns=["InternetService"] for small datasets partitioning is not recommended
    description="Telco customer features",
    tags={"source": "bronze", "format": "delta"}
)
```

```
from databricks.feature_engineering import FeatureEngineeringClient

fe = FeatureEngineeringClient()
```

Explore Feature Table with the UI

The screenshot shows the Databricks Catalog Explorer interface. On the left, the catalog tree displays a structure under 'menaf_gul_8': 'menaf_gul_8vg5_da' (expanded) > 'default' > 'telco_customer_features'. A yellow box highlights this path. On the right, the 'Feature Table Details' page for 'telco_customer_features' is shown. The title bar indicates the table is in the 'default' database of catalog 'menaf_gul_8vg5_da'. Below the title, it says 'Owner: menaf.gul@databricks.com' and 'Popularity: ... Size: 171KiB, 1 file Last Updated: 13 minutes ago'. A 'Create' button is also present. The 'Columns' tab is selected, showing the following schema:

Column	Type	Comment	Tags
customerID	string		
gender	string		
SeniorCitizen	int		
Partner	string		
Dependents	string		
tenure	int		
PhoneService	string		
MultipleLines	string		

Load Feature Table

- We can also look at the metadata of the feature store via the FeatureStore client by using `get_table()`.
- As feature table is a Delta table we can load it with Spark as normally we do for other tables.


```
ft = fe.get_table(name=table_name)
print(f"Feature Table description: {ft.description}")
print(ft.features)

Feature Table description: Telco customer features
['customerID', 'gender', 'SeniorCitizen', 'Partner', 'Dependents', 'tenure', 'PhoneService', 'MultipleLines', 'InternetService', 'OnlineSecurity', 'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV', 'StreamingMovies', 'Contract', 'PaperlessBilling', 'PaymentMethod', 'MonthlyCharges', 'TotalCharges']

display(fe.read_table(name=table_name))

Table
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 1   | 2   | 3   | 4   | 5   | 6   | 7   | 8   | 9   | 10  | 11  |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 1   | 2   | 3   | 4   | 5   | 6   | 7   | 8   | 9   | 10  | 11  |
| 7590-VHVEG | Female | 0 | Yes | No | 1 | No | No phone service | DSL |
| 5575-GNVED | Male | 0 | No | No | 34 | Yes | No | DSL |
| 3668-QPYBK | Male | 0 | No | No | 2 | Yes | No | DSL |
| 7795-CFOCW | Male | 0 | No | No | 45 | No | No phone service | DSL |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
```

Update Feature Table



- In some cases we might need to update an existing feature table by adding new features or deleting existing features. In this section, we will show to make these type of changes.
- Add a New Feature
 - To illustrate adding a new feature, let's redefine an existing one. In this case, we'll transform the tenure column by categorizing it into three groups: short, mid, and long, representing different tenure durations.
 - Then we will write the dataset back to the feature table. The important parameter is the mode parameter, which we should set to "**merge**".

Update Feature Table

pyspark.sql.DataFrame.withColumn

`DataFrame.withColumn(colName: str, col: pyspark.sql.Column) > pyspark.sql.dataframe.DataFrame`

Returns a new `DataFrame` by adding a column or re-

The column expression must be an expression over
`DataFrame` will raise an error.

```
>>> df = spark.createDataFrame([(2, "Alice"), (5,
>>> df.withColumn('age2', df.age + 2).show()
+---+---+
|age| name|age2|
+---+---+
| 2|Alice|  4|
| 5| Bob|  7|
+---+---+
```

```
from pyspark.sql.functions import when
telco_df_updated = telco_df.withColumn("tenure_group",
    when((telco_df.tenure >= 0) & (telco_df.tenure <= 25), "short")
    .when((telco_df.tenure > 25) & (telco_df.tenure <= 50), "mid")
    .when((telco_df.tenure > 50) & (telco_df.tenure <= 75), "long")
    .otherwise("invalid")
)
```

```
fe.write_table(
    name=table_name,
    df=telco_df_updated.select("customerID", "tenure_group"), # primary_key and column to add
    mode="merge"
)
```

Delete Existing Feature

- To remove a feature column from the table you can just drop the column. Let's drop the original tenure column.
- We need to set Delta read and write protocol version manually to support column mapping.
- Databricks supports column mapping for Delta Lake tables, which enables metadata-only changes to mark columns as deleted or renamed without rewriting data files.
- It also allows users to name Delta table columns using characters that are not allowed by Parquet, such as spaces, so that users can directly

ingest data into Delta tables due to the limitations of Parquet. This is particularly useful for legacy systems that have already been converted to Parquet but still need to interact with Delta tables.

27

```
%sql  
ALTER TABLE telco_customer_features SET TBLPROPERTIES ('delta.columnMapping.mode' = 'name', 'delta.minReaderVersion' = '2', 'delta.minWriterVersion' = '5');  
ALTER TABLE telco_customer_features DROP COLUMNS (tenure)
```

OK



AI Model Serving

ML workflow using feature engineering



Write code to convert raw data into features and create a Spark DataFrame containing the desired features.



[Create a Delta table in Unity Catalog](#) that has a primary key.



Train and log a model using the feature table. When you do this, the model stores the specifications of features used for training.



When the model is used for inference, it automatically joins features from the appropriate feature tables.




Register model in [Model Registry](#).

<https://learn.microsoft.com/en-us/azure/databricks/machine-learning/feature-store/>

Batch Use Case

For batch use cases, the model automatically retrieves the features it needs from Feature Store.



Real-time Use Case



For real-time serving use cases, publish the features to an [online table](#).

Third-party online stores are also supported.



An **online table** is a read-only copy of a Delta Table that is stored in row-oriented format optimized for online access.

Online tables are fully serverless tables that auto-scale throughput capacity with the request load and provide low latency and high throughput access to data of any scale.

Online tables are designed to work with:

- Mosaic AI Model Serving,
- Feature Serving, and
- Retrieval-augmented generation (RAG) applications where they are used for fast data lookups.

Online Table (Databricks UI)

The screenshot shows the Databricks Catalog Explorer interface. On the left, there's a sidebar with a search bar and a tree view of catalogs and tables. The main area shows the details for a table named 'example' in the 'docs.default' catalog. The table has two columns: 'id' (bigint) and 'random' (double). A context menu is open on the right, listing options like 'Query', 'Notebook', 'Dashboard', 'Vector search index', and 'Online table'. The 'Online table' option is highlighted.

Column	Type	Comment	Tags
id	bigint		
random	double		

<https://learn.microsoft.com/en-us/azure/databricks/machine-learning/feature-store/online-tables#api-sdk>

Online Table (Databricks UI)

Create online table

Name
docs. default

Primary Key

Timeseries Key ⓘ

Sync mode [sync modes explained ↗](#)

Snapshot Triggered Continuous

Online Table (Databricks UI)

The screenshot shows the Databricks Catalog interface. On the left, the Catalog sidebar lists databases: 'docs' (expanded), 'default', 'docs_user' (expanded), 'online_wine' (selected and highlighted in blue), 'wine' (under 'docs'), and 'information_schema'. The main panel displays the details for the 'online_wine' table under the 'docs_user' database. The title is 'Catalogs > docs > docs_user > docs.docs_user.online_wine'. The table is described as an 'Online view of Delta table'. The 'Overview' tab is selected, showing the following information:

Setting	Value
Online table status	Online
Source table	docs.docs_user.wine
Primary key(s)	["id"]
Sync schedule	Snapshot
Data ingest	
Pipeline id	[redacted]
Update status	Completed
Last processed timestamp	Dec 05, 2023, 02:18 PM (8 minutes ago)

A 'Sync now' button is located at the bottom right of the table details area.


Sync Mode

Policy	Description
Snapshot	The pipeline runs once to take a snapshot of the source table and copy it to the online table. Subsequent changes to the source table are automatically reflected in the online table by taking a new snapshot of the source and creating a new copy. The content of the online table is updated atomically.
Triggered	The pipeline runs once to create an initial snapshot copy of the source table in the online table. Unlike the Snapshot sync mode, when the online table is refreshed, only changes since the last pipeline execution are retrieved and applied to the online table. The incremental refresh can be manually triggered or automatically triggered according to a schedule.
Continuous	The pipeline runs continuously. Subsequent changes to the source table are incrementally applied to the online table in real time streaming mode. No manual refresh is necessary.

ⓘ Note

To support **Triggered** or **Continuous** sync mode, the source table must have [Change data feed](#) enabled.

3rd Party Online Tables



<https://learn.microsoft.com/en-us/azure/databricks/machine-learning/feature-store/publish-features>

3rd Party Online Tables

```
Python

import datetime
from databricks.feature_engineering.online_store_spec import AzureMySqlSpec
# or databricks.feature_store.online_store_spec for Workspace Feature Store
online_store = AzureMySqlSpec(
    hostname='<hostname>',
    port='<port>',
    read_secret_prefix='<read-scope>/<prefix>',
    write_secret_prefix='<write-scope>/<prefix>'
)

fs.publish_table(
    name='recommender_system.customer_features',
    online_store=online_store,
    filter_condition=f"_dt = '{str(datetime.date.today())}'",
    mode='merge'
)
```

```
Python

import datetime
from databricks.feature_engineering.online_store_spec import AzureCosmosDBSpec
# or databricks.feature_store.online_store_spec for Workspace Feature Store
online_store = AzureCosmosDBSpec(
    account_uri='<account-uri>',
    read_secret_prefix='<read-scope>/<prefix>',
    write_secret_prefix='<write-scope>/<prefix>'
)

fe.publish_table( # or fs.publish_table for Workspace Feature Store
    name='ml.recommender_system.customer_features',
    online_store=online_store,
    filter_condition=f"_dt = '{str(datetime.date.today())}'",
    mode='merge'
)
```

Mosaic AI Model Serving

- Mosaic AI Model Serving provides a unified interface to deploy, govern, and query AI models for real-time and batch inference.
- Each model you serve is available as a REST API that you can integrate into your web or client application.
- Model Serving provides a highly available and low-latency service for deploying models.
- The service automatically scales up or down to meet demand changes, saving infrastructure costs while optimizing latency performance.
 - This functionality uses serverless compute.
- It also provides a single UI to manage all your models and their respective serving endpoints. You can also access models directly from SQL using [AI Functions](#) for easy integration into analytics workflows.

<https://docs.databricks.com/aws/en/machine-learning/model-serving/>


<https://docs.databricks.com/aws/en/large-language-models/ai-functions>

Why use Model Serving

- **Deploy and query any models:** Model Serving provides a unified interface that so you can manage all models in one location and query them with a single API, regardless of whether they are hosted on Databricks or externally.
 - This approach simplifies the process of experimenting with, customizing, and deploying models in production across various clouds and providers.
- **Securely customize models with your private data:** Built on a Data Intelligence Platform, Model Serving simplifies the integration of features and embeddings into models through native integration with the [Databricks Feature Store](#) and [Mosaic AI Vector Search](#).
 - For even more improved accuracy and contextual understanding, models can be fine-tuned with proprietary data and deployed effortlessly on Model Serving.

<https://docs.databricks.com/aws/en/machine-learning/model-serving/>

<https://learn.microsoft.com/en-us/azure/databricks/resources/feature-region-support#azure-model-serving>



Mosaic AI Gateway

- Mosaic AI Gateway is designed to streamline the usage and management of generative AI models and agents within an organization.
 - It is a centralized service that brings governance, monitoring, and production readiness to model serving endpoints.
 - It also allows you to run, secure, and govern AI traffic to democratize and accelerate AI adoption for your organization.

<https://learn.microsoft.com/en-us/azure/databricks/resources/feature-region-support#azure-model-serving>

AI Gateway Features

Feature	Definition
Permission and rate limiting	Control who has access and how much access.
Payload logging	Monitor and audit data being sent to model APIs using inference tables .
Usage tracking	Monitor operational usage on endpoints and associated costs using system tables .
AI Guardrails	Prevent unwanted and unsafe data in requests and responses. See AI Guardrails .
Fallbacks	Minimize production outages during and after deployment.
Traffic splitting	Load balance traffic across models.

<https://docs.databricks.com/aws/en/ai-gateway/>