



SQL Server Table Partitioning

© Copyright Microsoft Corporation. All rights reserved.

CS340 | SBASSE | LUMS

Table Partitioning

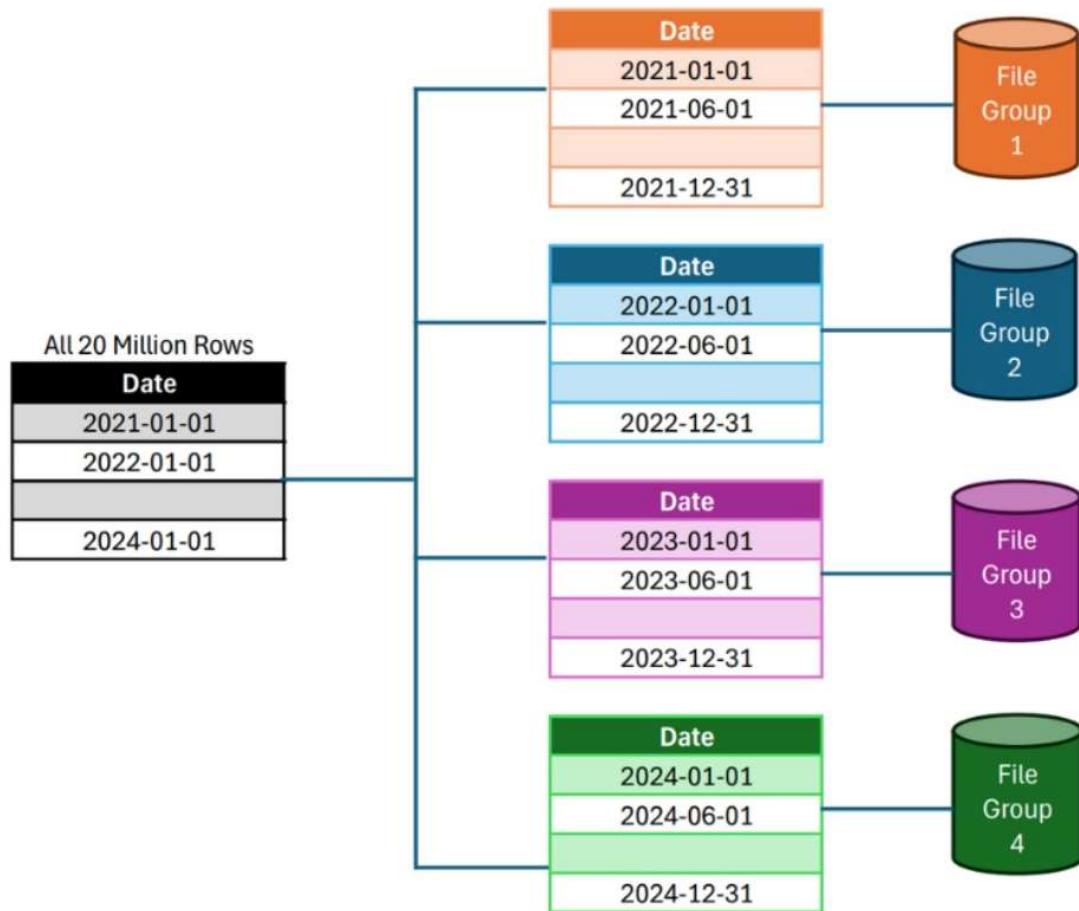
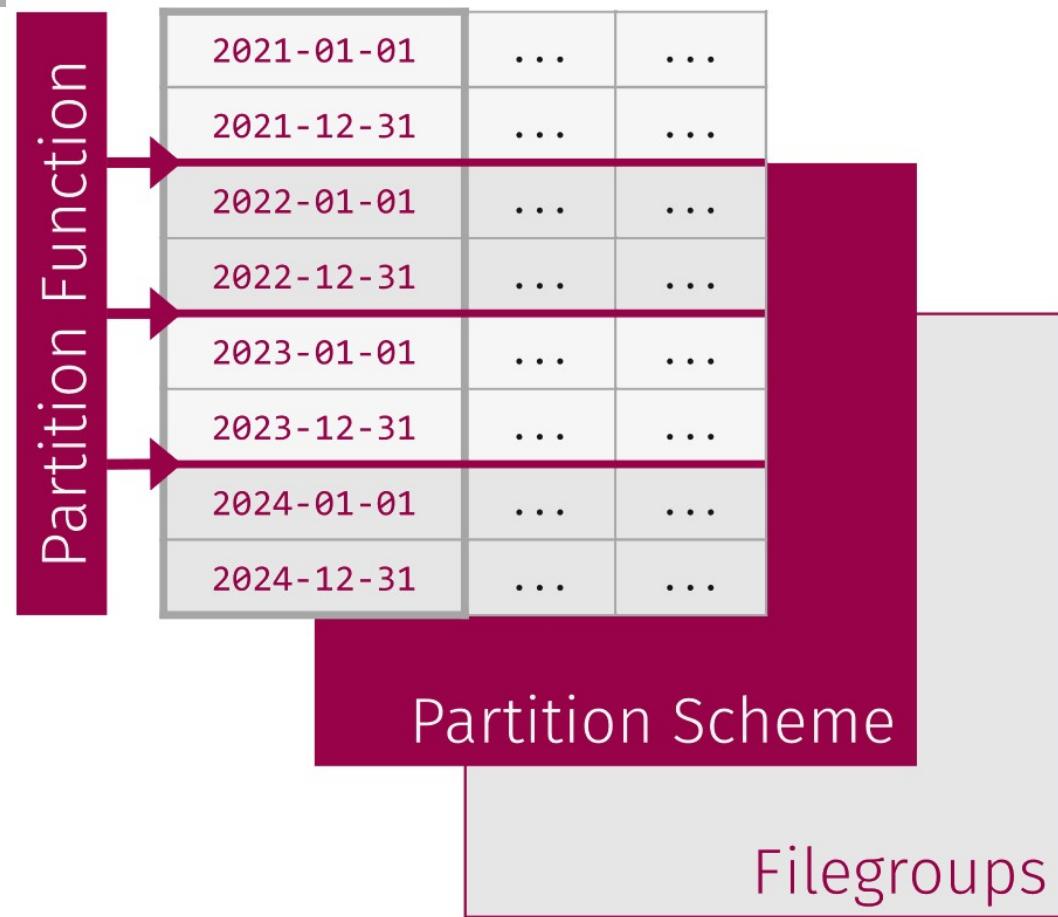
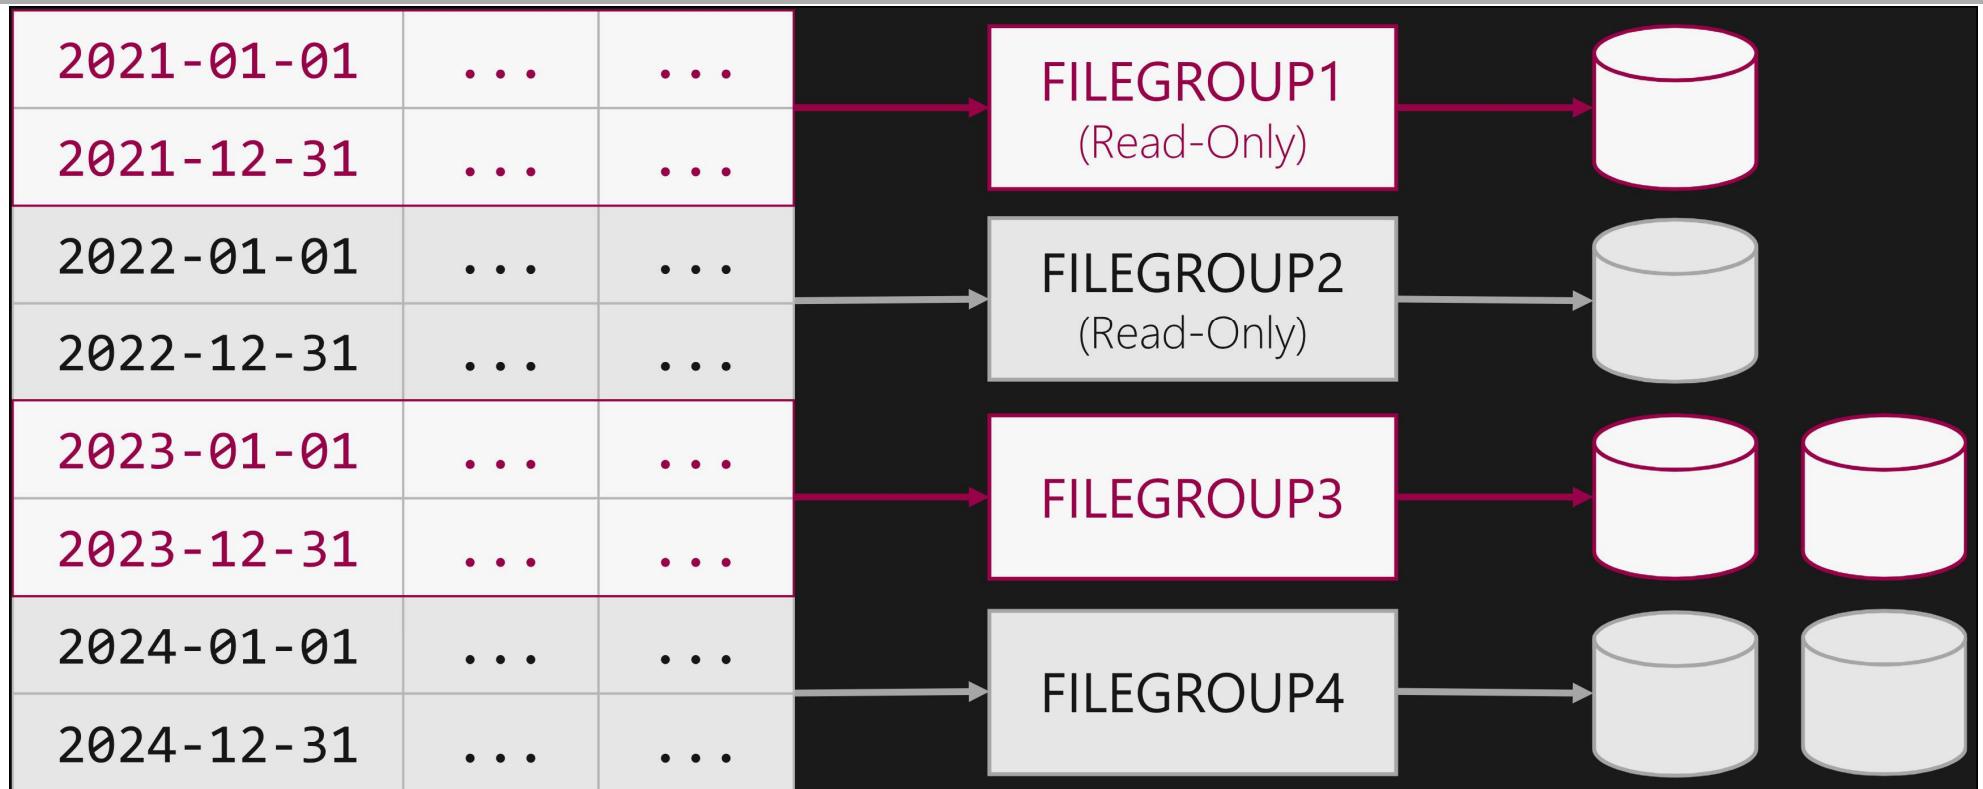
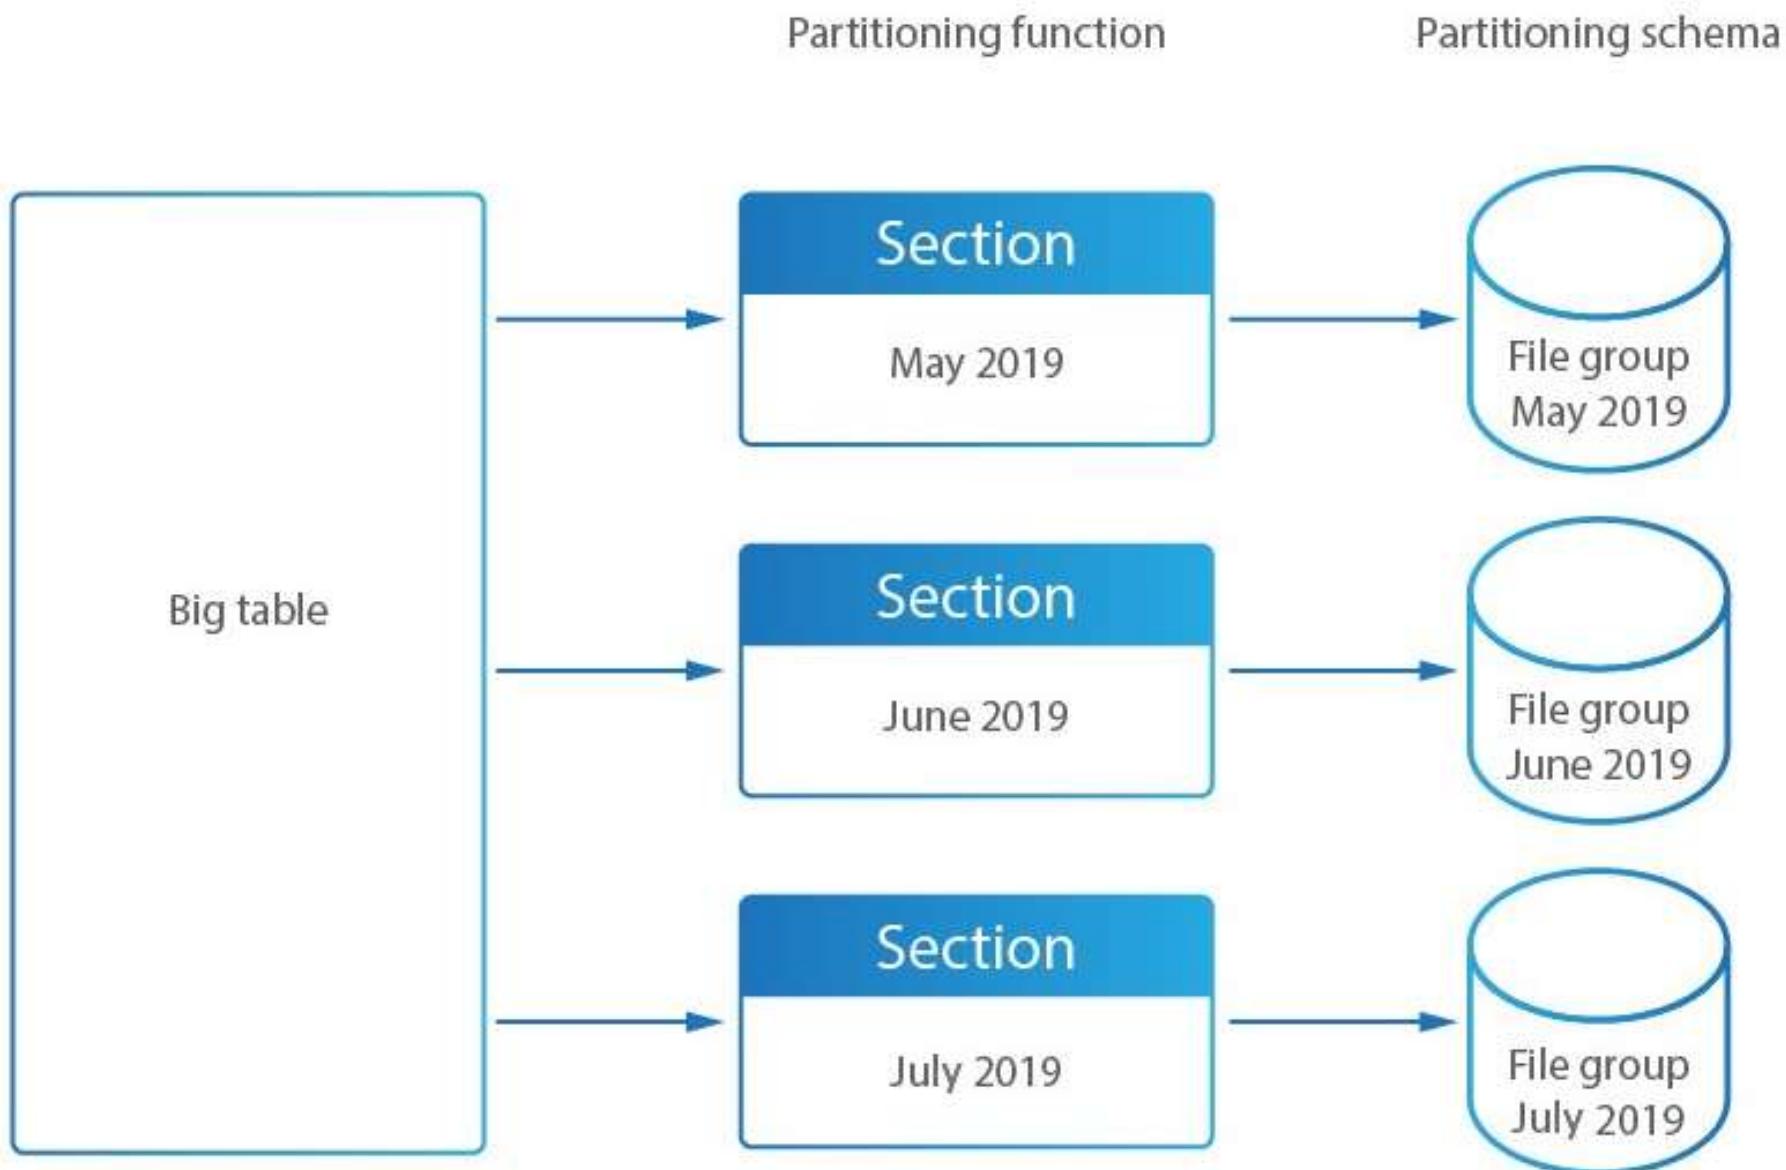


Table Partitioning



Partitioning Scheme

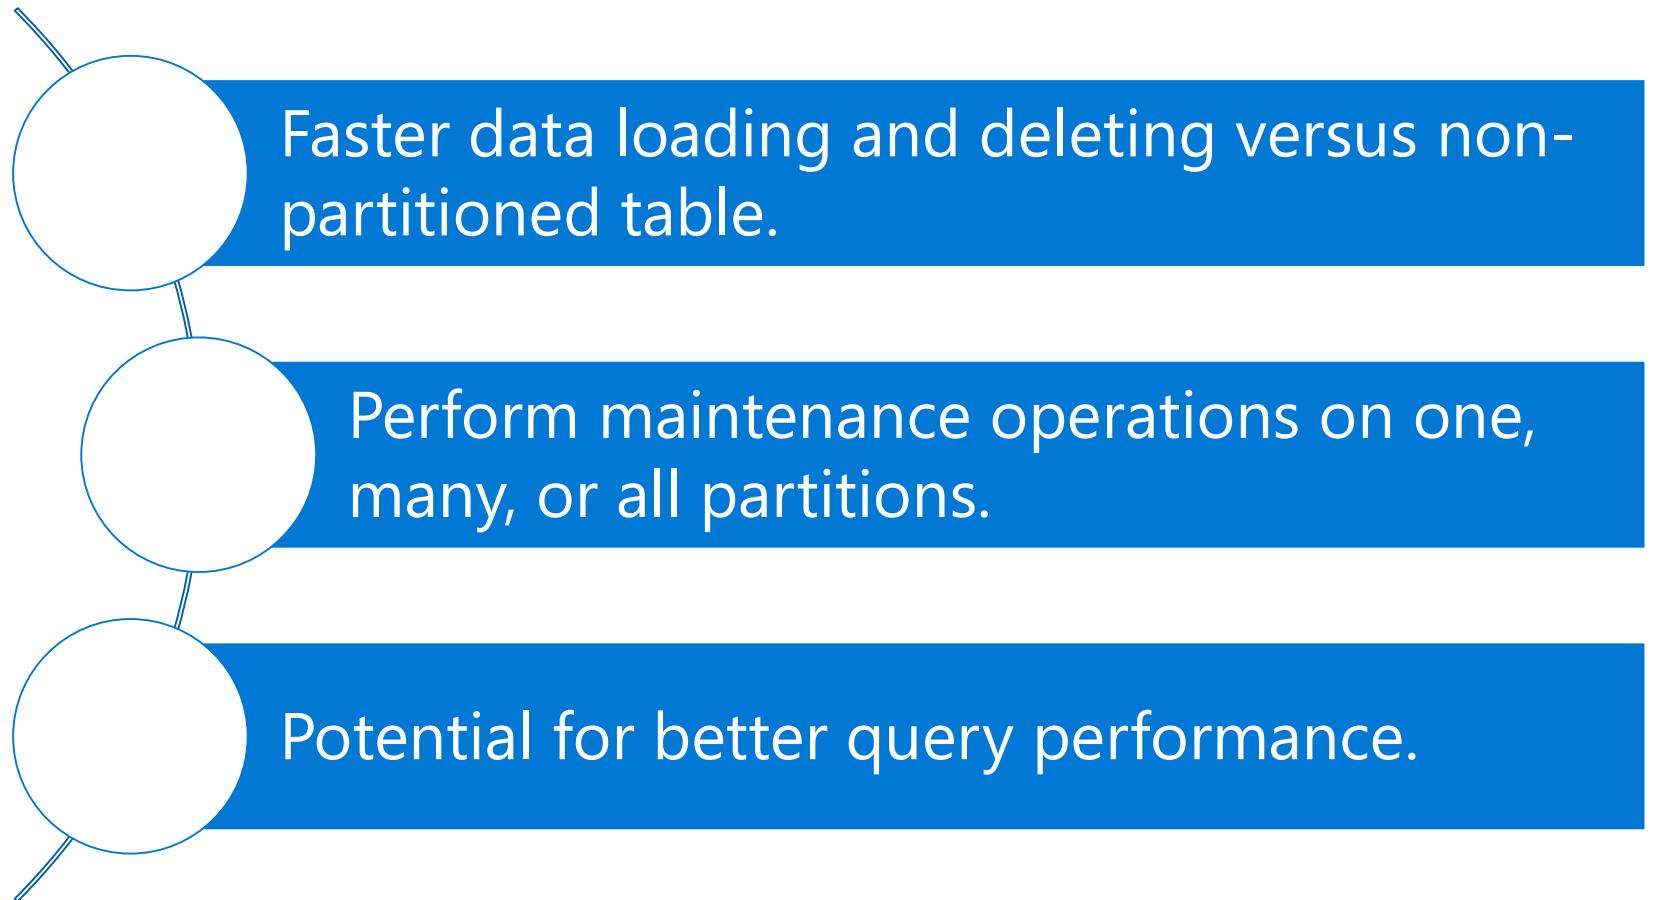




What is Table Partitioning

- Technique in SQL Server that allows you to split a large table into smaller, more manageable pieces (partitions) based on the values in a specific column
- Each partition behaves like a separate table internally
- Partitions are invisible to the user: Treated as one entity when queried.

Why use partitioning?

- 
- Faster data loading and deleting versus non-partitioned table.
 - Perform maintenance operations on one, many, or all partitions.
 - Potential for better query performance.

When do you choose partitioning?

- Large table causes maintenance or performance concerns
- When to consider:
 - Removing a large volume of data is slow or causes blocking
 - Loading a large volume of data is slow or causes blocking
 - Backup / maintenance exceeds maintenance window
- Partitioning for performance
 - Parallel operations on large tables
 - Prunes data easier, reduces I/O

Is it in all versions/editions?

- Prior to SQL Server 2016 SP1
 - No, Enterprise Edition only.
- As of SQL Server 2016 SP1+, yes.
- Up to 15,000 partitions
 - Prior to SQL Server 2012, limited to 1,000.



Table partitioning basics



Partitioning Building Blocks

1. Partitioning column
2. Partitioning function
3. Partitioning scheme

1. Partitioning Column

- Must be a [computed] column on the table
- A valid data type
 - (not ntext, text, image, xml, varchar(max), nvarchar(max), or varbinary(max))
- If clustered, column must be PK or the clustered index.

- If partition column is not PK:
 - Avoid Nullable columns in partition colum
 - NULL partition column will reside in leftmost partition
- Partition column to divide table
 - Should be relatively balanced
- A common filtering column
 - Enables partition elimination in query processing

2. Partition Function

- Defines how the rows will be mapped to the partitions
- Based on the partitioning column

```
CREATE PARTITION FUNCTION OrderDatePF (datetime)
AS RANGE RIGHT FOR VALUES
('2018/01/01', '2019/01/01', '2020/01/01');
```

RANGE RIGHT means that each boundary value belongs to the partition to its right.

Function – OrderDatePF

Partition1

OrderDate
 $< 2018/01/01$

Partition2

OrderDate
 $\geq 2018/01/01$
AND
 $< 2019/01/01$

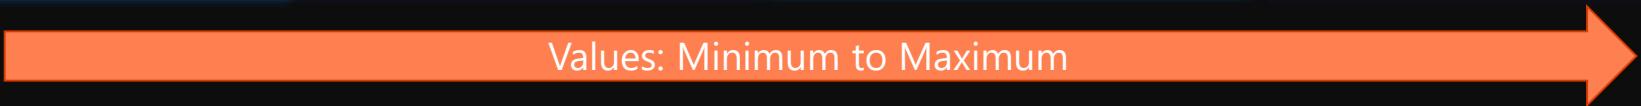
Partition3

OrderDate
 $\geq 2019/01/01$
AND
 $< 2020/01/01$

Partition4

OrderDate
 $\geq 2020/01/01$

Values: Minimum to Maximum



Range Types – RIGHT or LEFT

- As part of the function, you need to determine if your partition will range right or range left.
- Range right - "begins at" - lower boundary
- Range left - "ends at" - upper boundary

```
CREATE PARTITION FUNCTION OrderDatePF (datetime)
AS RANGE RIGHT FOR VALUES
('2018/01/01', '2019/01/01', '2020/01/01');
```

Note: RANGE RIGHT would require the full timestamp to keep date values of the same year in the same partition. Example: '2019/01/01 23:59:59.997'

Range Type Example – Unpartitioned Table



RANGE RIGHT

```
CREATE PARTITION FUNCTION PF1_Right (int)  
AS RANGE RIGHT FOR VALUES (10);
```

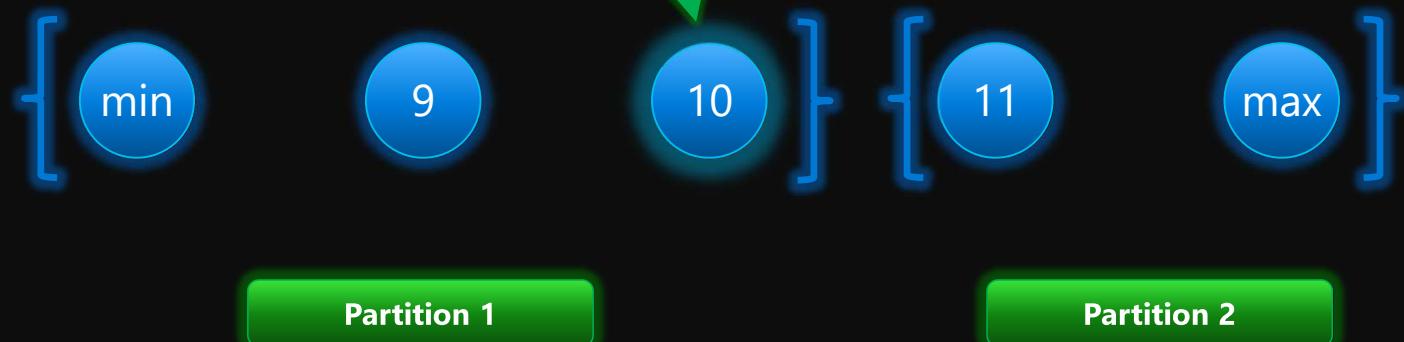


Partition 1

Partition 2

RANGE LEFT

```
CREATE PARTITION FUNCTION PF1_Left (int)  
AS RANGE LEFT FOR VALUES (10);
```





Partitioning Scheme



Partitioning Scheme

A **partitioning scheme** defines **how data is distributed across partitions** based on a **partition function**.

- The partition function, as defined earlier, maps rows to partitions based on the values in a specified column (often a date or ID).
- Then partitioning scheme maps those partitions to filegroups.

Benefits of Partitioning:

- **Improved Query Performance:** Queries can target specific partitions.
- **Efficient Maintenance:** You can switch, truncate, or rebuild partitions individually.
- **Better Storage Management:** Distribute data across multiple filegroups/disks.

Partitioning Scheme to Map Each Partition to different
filegroup

```
CREATE PARTITION SCHEME YearPS
AS PARTITION OrderDatePF
TO (FG1, FG2, FG3, FG4);
```

```
CREATE TABLE Orders
(OrderId INT, OrderDate DATETIME, Col3,etc...)
ON YearPS(OrderDate);
```

Partitioning Scheme: Example

Imagine a table storing sales data for multiple years. You can partition it by year:

```
CREATE PARTITION FUNCTION pfYear(INT)
AS
RANGE LEFT FOR VALUES (2019,2020,2021,2022);
```

The above maps each year to a different filegroup and creates a table using this scheme. Queries for a specific year will only scan the relevant partition.

Partitioning Scheme

